

Algorithms & Data Structures

Lecture 08 Binary Search Trees

Giovanni Bacci
giovbacci@cs.aau.dk

Outline

- Binary Search Trees
- Operations on Binary Search Trees
- Red-Back Trees
- Operations on Red-Back Trees

Intended Learning Goals

KNOWLEDGE

- Mathematical reasoning on concepts such as recursion, induction, concrete and abstract computational complexity
- Data structures, algorithm principles e.g., search trees, hash tables, dynamic programming, divide-and-conquer
- Graphs and graph algorithms e.g., graph exploration, shortest path, strongly connected components.

SKILLS

- Determine abstract complexity for specific algorithms
- Perform complexity and correctness analysis for simple algorithms
- Select and apply appropriate algorithms for standard tasks

COMPETENCES

- Ability to face a non-standard programming assignment
- Develop algorithms and data structures for solving specific tasks
- Analyse developed algorithms

Why Binary Search Trees

- Binary search trees support many dynamic-set operations, including Search, Minimum, Maximum, Predecessor, Successor, Insert, and Delete
- We can use search trees both as dictionary and as priority queue.
- Many basic operations on binary search trees take time proportional to the height of the tree. If the tree is balanced this corresponds to $\Theta(\lg n)$

Binary Search Trees

Binary Search Trees

- BSTs can be represented by means of linked data structures where each node has
 - A key (and optional satellite data)
 - Links to its parent node, left child, and right child
- The root node is the only node in the tree whose parent is Nil
- Keys in the tree satisfy the binary-search-tree property

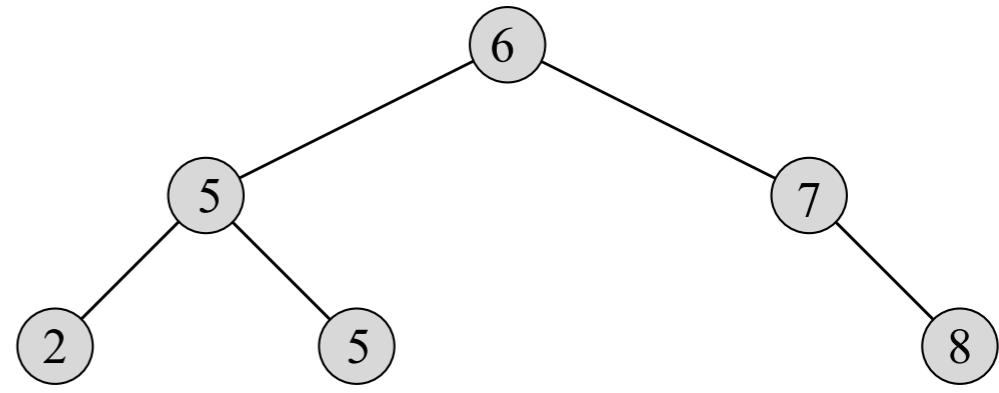
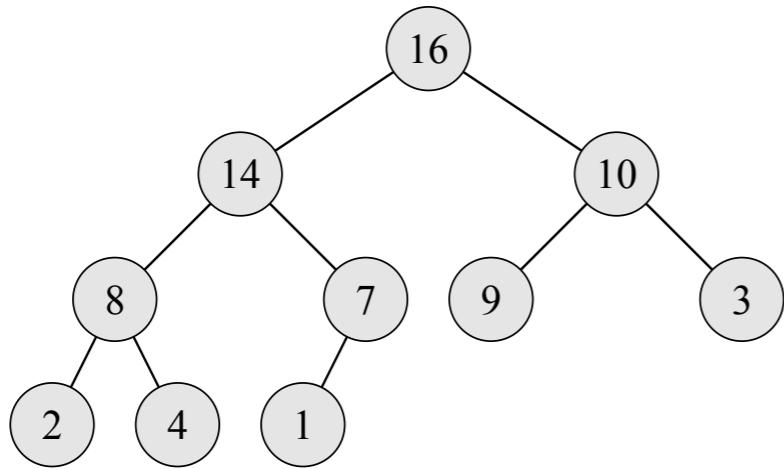
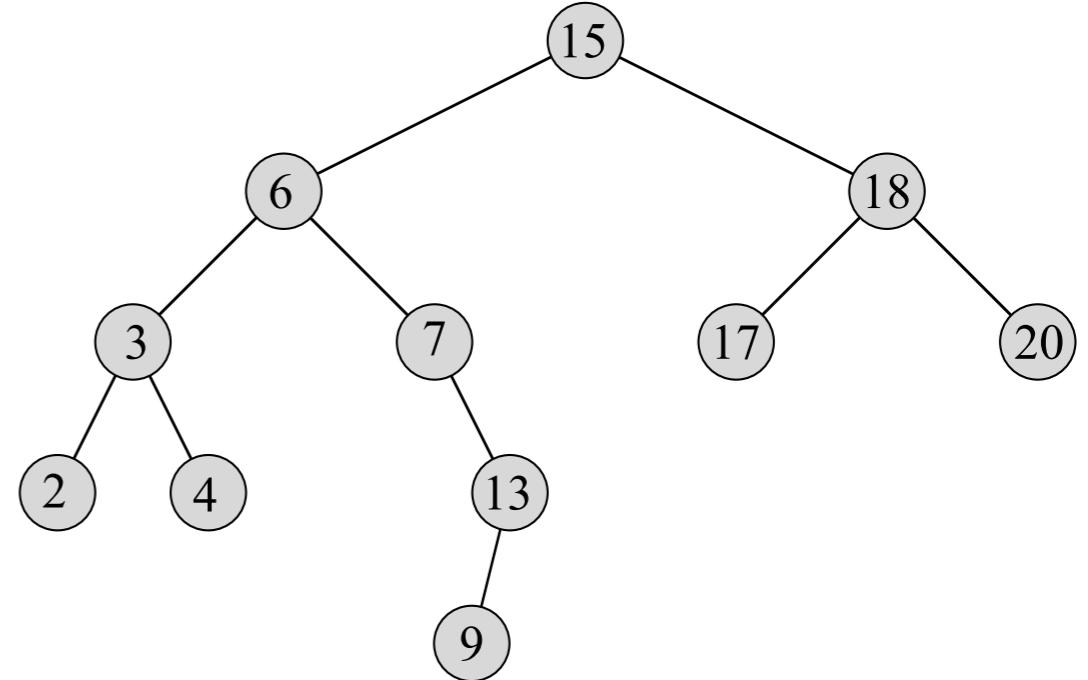
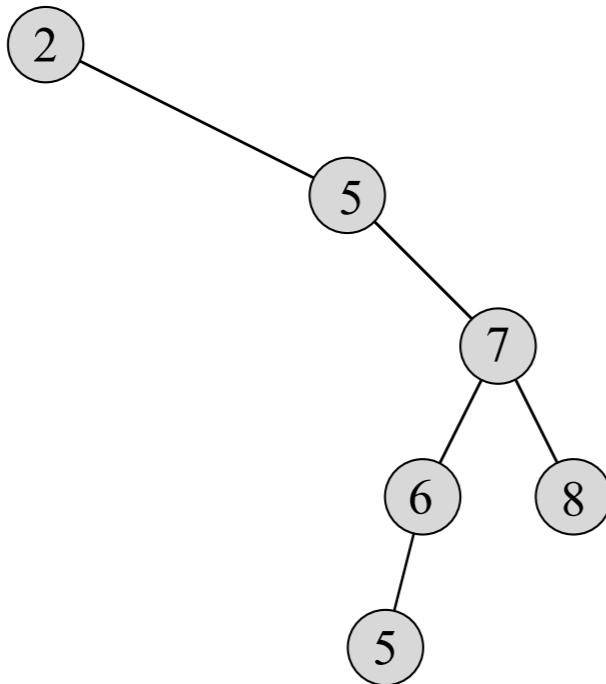
Binary-search-tree property:

For any node node x in the tree, the following hold

- If y is a node in the left subtree of x then $y.key \leq x.key$
- If z is a node in the right subtree of x then $x.key \leq z.key$

Quiz

Which of the following is a BST?



Inorder Tree Walk

- The binary-search-tree property allows one to print out all the keys in a BST in sorted order by a simple inorder visit of the tree, called **inorder tree walk**

INORDER-TREE-WALK(x)

```
1  if  $x \neq \text{NIL}$ 
2      INORDER-TREE-WALK( $x.\text{left}$ )
3      print  $x.\text{key}$ 
4      INORDER-TREE-WALK( $x.\text{right}$ )
```

First recursively visit the left child, then print the node x , and recursively visit the right child

- Recall from Exercise Session 05 that this way of visiting the tree takes $\Theta(n)$ time for an n -node binary tree (see also Theorem 12.1 CLRS)

Other Visiting Orders

Similarly to the **inorder tree walk**, we can visit the tree using

- **Preorder tree walk:** prints the root before making the recursive calls to the left and right children
- **Postorder tree walk:** prints the root after returning from the recursive calls to the left and right children
- Both procedures take $\Theta(n)$.

Quiz

Can you think of a procedure that prints all the elements of a binary search tree in reversed-sorted order?

Answer

Can you think of a procedure that prints all the elements of a binary search tree in reversed-sorted order?

INORDER-TREE-WALK(x)

- 1 **if** $x \neq \text{NIL}$
- 2 INORDER-TREE-WALK($x.\text{left}$)
- 3 print $x.\text{key}$
- 4 INORDER-TREE-WALK($x.\text{right}$)



Just switch these two recursive calls

Searching in a BST

- Given a pointer to the root of the tree and a key k , **TREE-SEARCH** returns a pointer to a node with key k if it exists; otherwise it returns **NIL**
- Exploit the BST property** to choose the appropriate subtree
- Worst-case running time** $O(h)$ where h is the height of the tree rooted at x

TREE-SEARCH(x, k)

```
1  if  $x == \text{NIL}$  or  $k == x.key$ 
2      return  $x$ 
3  if  $k < x.key$ 
4      return TREE-SEARCH( $x.left, k$ )
5  else return TREE-SEARCH( $x.right, k$ )
```

ITERATIVE-TREE-SEARCH(x, k)

```
1  while  $x \neq \text{NIL}$  and  $k \neq x.key$ 
2      if  $k < x.key$ 
3           $x = x.left$ 
4      else  $x = x.right$ 
5  return  $x$ 
```

Minimum and Maximum

- Given a pointer x to the root of the tree and a key k , **MINIMUM** and **MAXIMUM** resp. return a pointer to a node with minimal and maximal key k
- Exploit the **BST property to avoid comparisons**
- Worst-case running time** $O(h)$ where h is the height of the tree rooted at x

TREE-MINIMUM(x)

```
1 while  $x.left \neq \text{NIL}$ 
2      $x = x.left$ 
3 return  $x$ 
```

TREE-MAXIMUM(x)

```
1 while  $x.right \neq \text{NIL}$ 
2      $x = x.right$ 
3 return  $x$ 
```

(The same procedures can be equivalently implemented using recursion)

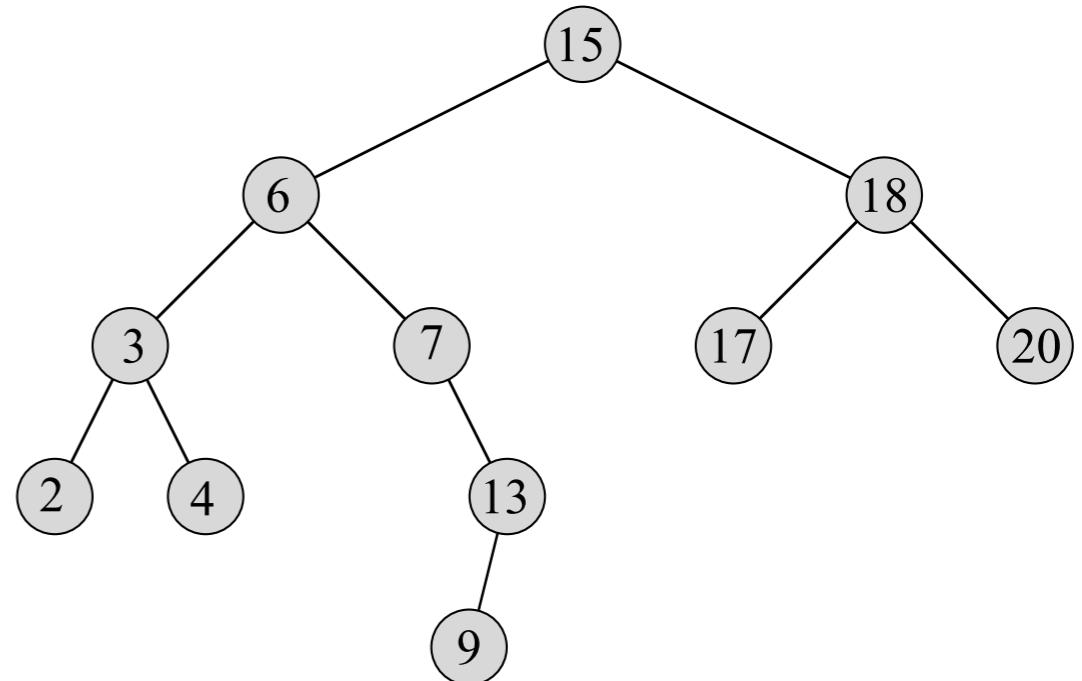
Successor

- We assume that all nodes in the tree are distinct.
- Given a node x in a BST, the procedure Tree-Successor(x) returns the node having the smallest key greater than $x.key$ if it exists; otherwise it returns Nil

Quiz

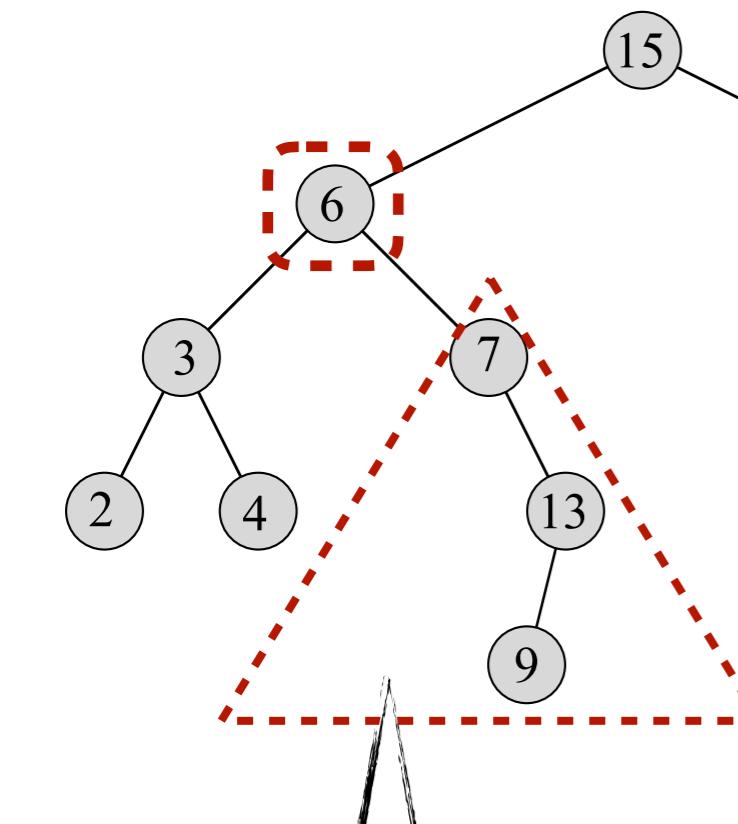
Consider the BST to the right

- Who is the successor of 6?
- And the successor of 13?
- Can we use the BST property to avoid comparisons?

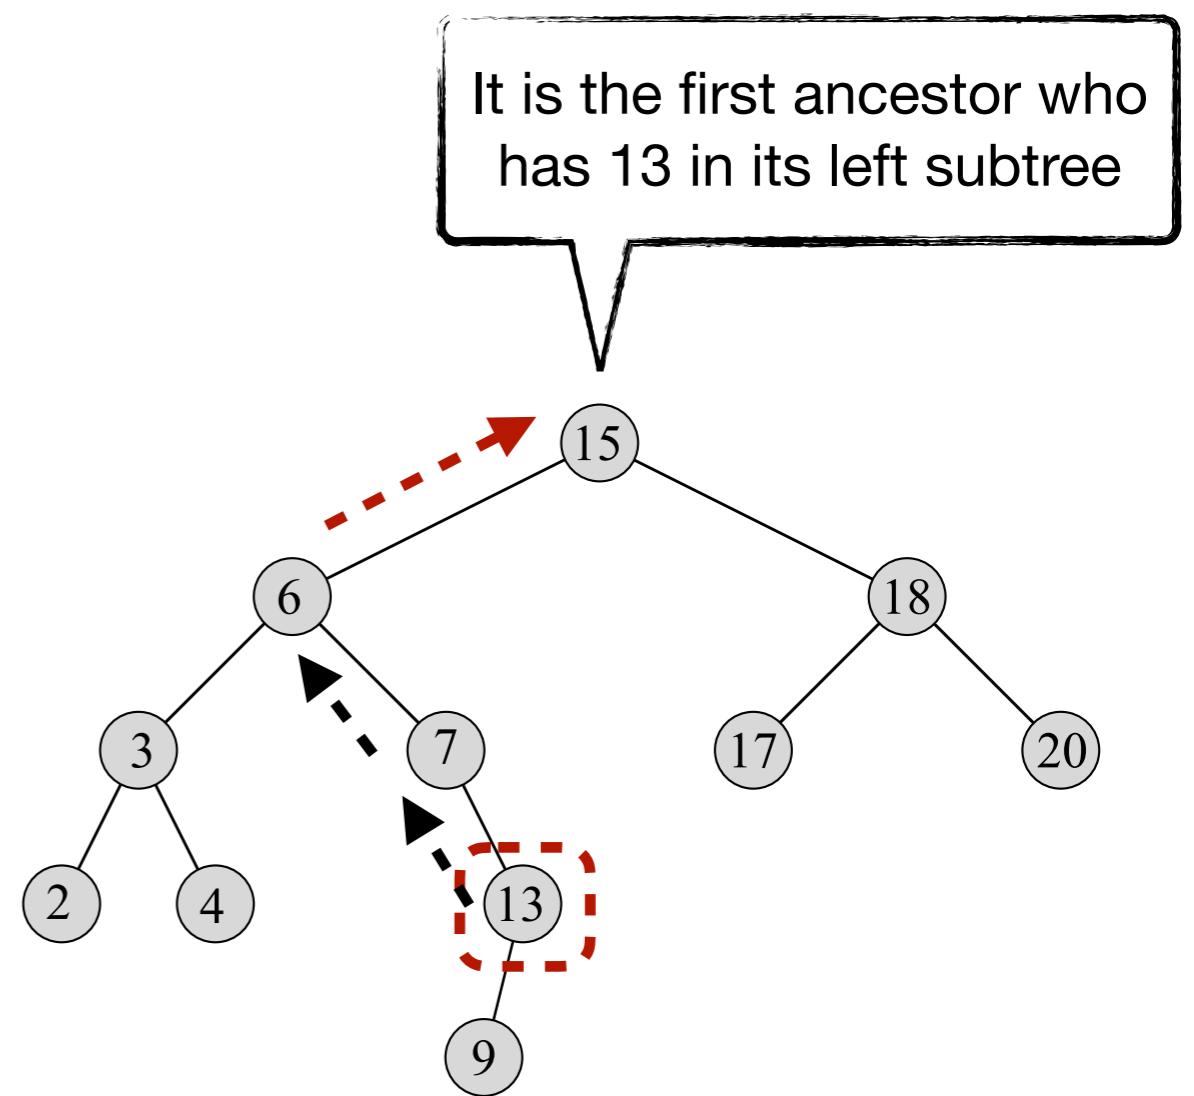


Answer

We can exploit the BST property for both cases!



It is the Minimum of the
subtree rooted as $x.right$



It is the first ancestor who
has 13 in its left subtree

Successor

- We assume that all nodes in the tree are distinct.
- Given a node x in a BST, the procedure Tree-Successor(x) returns the node having the smallest key greater than $x.key$ if it exists; otherwise it returns Nil

TREE-SUCCESSOR(x)

```
1  if  $x.right \neq \text{NIL}$ 
2      return TREE-MINIMUM( $x.right$ )
3
4  while  $y \neq \text{NIL}$  and  $x == y.right$ 
5       $x = y$ 
6       $y = y.p$ 
7  return  $y$ 
```

Case 1: minimum element
in the right subtree

Case 2: move to the first
ancestor who has x in the
left subtree, by climbing up
the parent pointers

Tree-Successor: run-time

- The running time of successor and predecessor operations is $\Theta(h)$ where h is the height of the whole BST

TREE-SUCCESSOR(x)

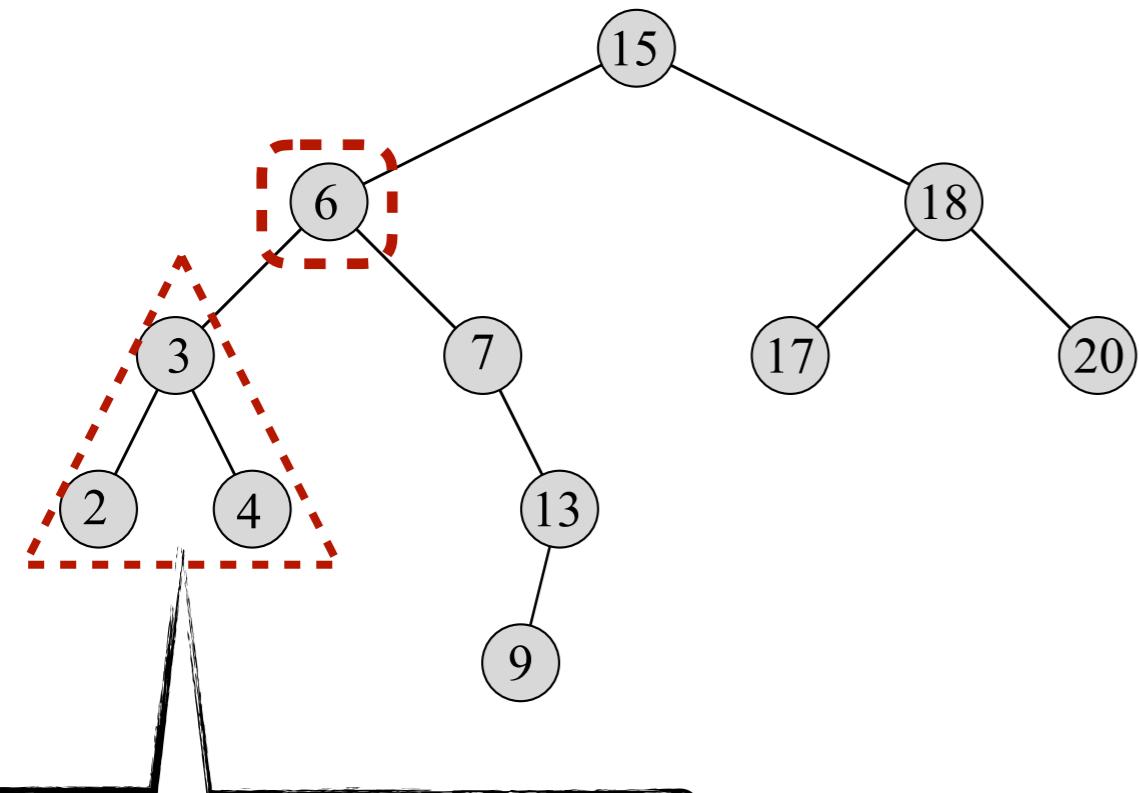
```
1  if  $x.right \neq \text{NIL}$ 
2      return TREE-MINIMUM( $x.right$ )
3   $y = x.p$ 
4  while  $y \neq \text{NIL}$  and  $x == y.right$ 
5       $x = y$ 
6       $y = y.p$ 
7  return  $y$ 
```

Case 1: we have seen
before $\Theta(h)$

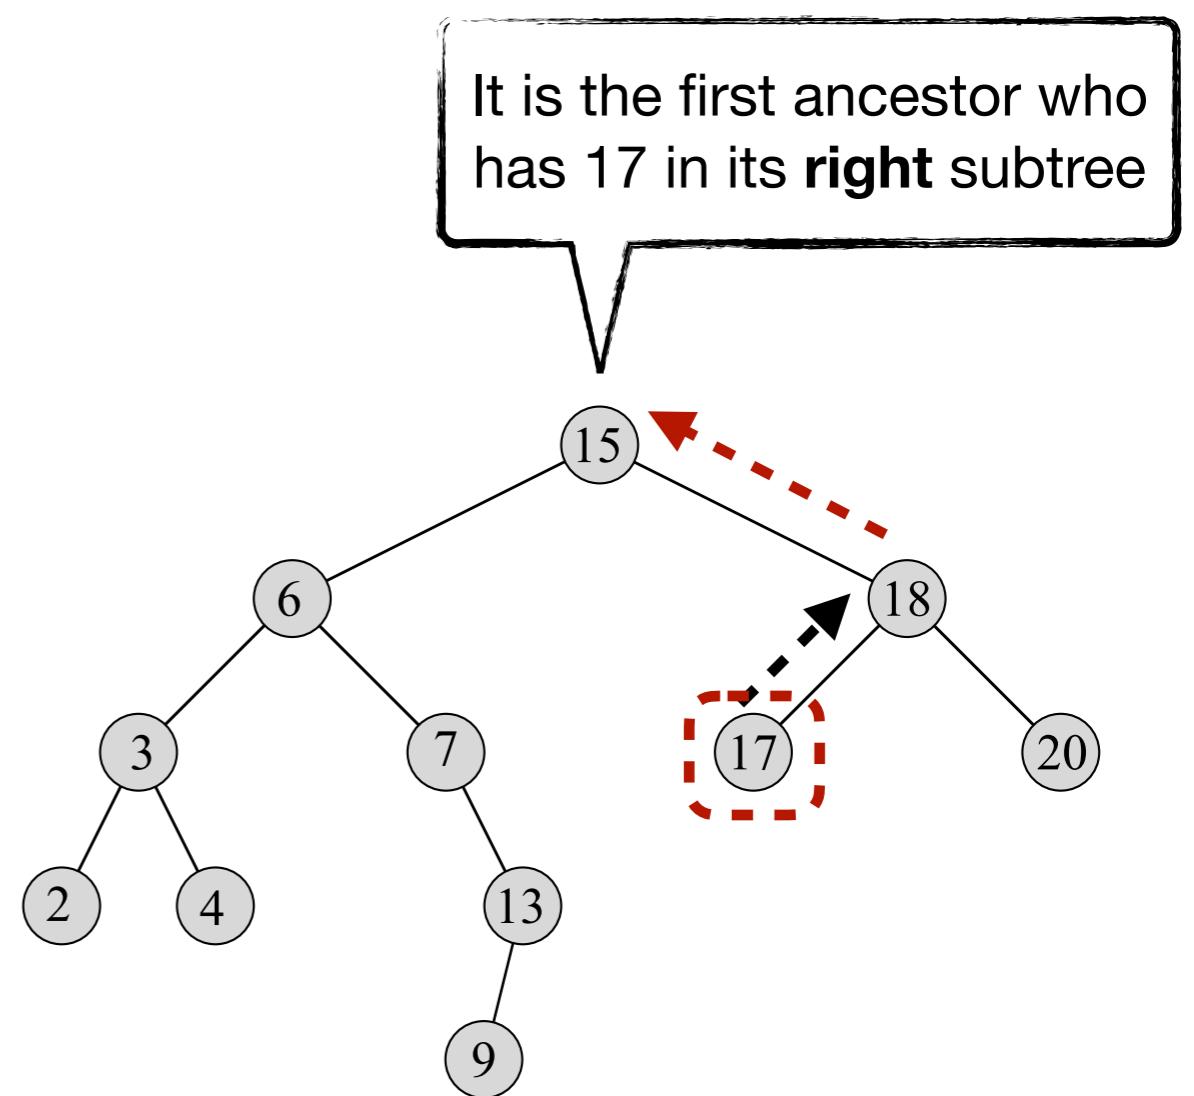
Case 2: it visits a simple path
which stops in the worst-case
at the root.
The length of the path is $\Theta(h)$

What about Predecessor?

The procedure is analogous to tree-successor, and has same running-time $\Theta(h)$



It is the **Maximum** of the subtree rooted as $x.left$



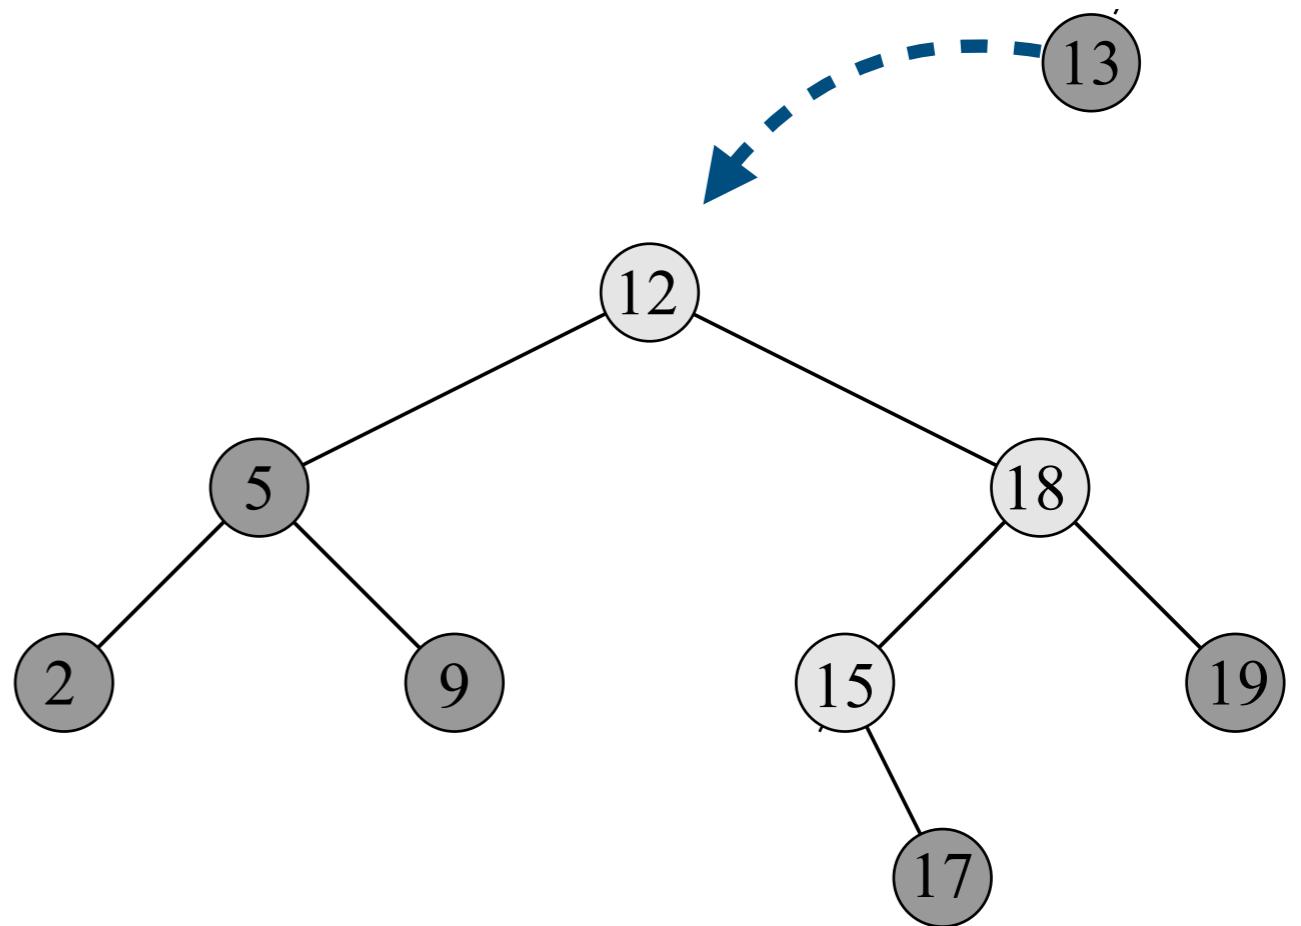
It is the first ancestor who has 17 in its **right** subtree

Insertion

The `TREE-INSERT` procedure takes a BST T and a node z and update the links so that z becomes a node of T while preserving the BST property

Basic Idea:

- Place the node as a new leaf of the tree
- Ensure the BST property holds by choosing an appropriate path starting from the root

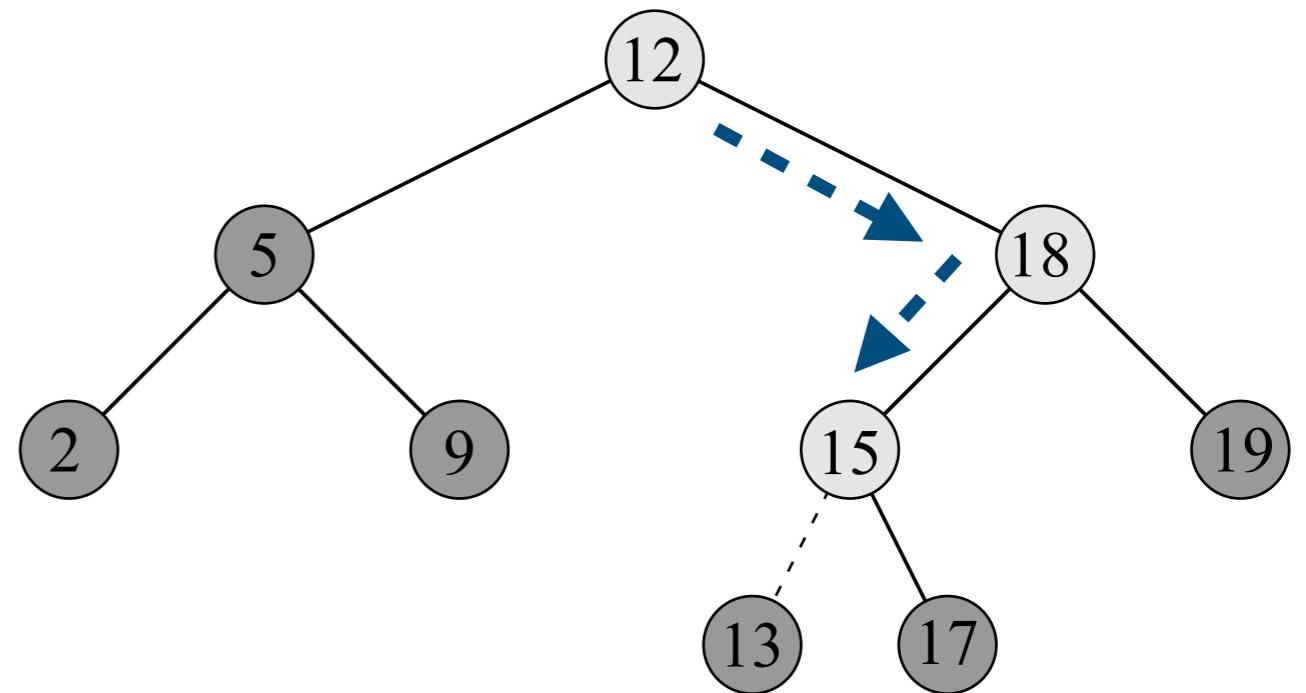


Insertion

The `TREE-INSERT` procedure takes a BST T and a node z and updates the pointers so that z becomes a node of T while preserving the BST property

Basic Idea:

- Place the node as a new leaf of the tree
- Ensure the BST property holds by choosing an appropriate path starting from the root



Tree-Insert: pseudocode

The **TREE-INSERT** procedure takes a BST T and a node z and update the links so that z becomes a node of T while preserving the BST property

TREE-INSERT(T, z)

```
1   $y = \text{NIL}$ 
2   $x = T.\text{root}$ 
3  while  $x \neq \text{NIL}$ 
4       $y = x$ 
5      if  $z.\text{key} < x.\text{key}$ 
6           $x = x.\text{left}$ 
7      else  $x = x.\text{right}$ 
8   $z.p = y$ 
9  if  $y == \text{NIL}$ 
10      $T.\text{root} = z$ 
11  elseif  $z.\text{key} < y.\text{key}$ 
12       $y.\text{left} = z$ 
13  else  $y.\text{right} = z$ 
```

x is the current node
and y its parent

Move x and y down in the tree
comparing $x.\text{key}$ with $z.\text{key}$

At this point x is Nil and z can replace it

Case 1: T was empty. Then, z becomes the root

Case 2: place z in the appropriate
subtree by comparing $z.\text{key}$ with $y.\text{key}$

Tree-Insert: run-time

The **TREE-INSERT** procedure takes a BST T and a node z and update the links so that z becomes a node of T while preserving the BST property

TREE-INSERT(T, z)

```
1   $y = \text{NIL}$ 
2   $x = T.\text{root}$ 
3  while  $x \neq \text{NIL}$ 
4       $y = x$ 
5      if  $z.\text{key} < x.\text{key}$ 
6           $x = x.\text{left}$ 
7      else  $x = x.\text{right}$ 
8
9       $z.p = y$ 
10     if  $y == \text{NIL}$ 
11          $T.\text{root} = z$ 
12     elseif  $z.\text{key} < y.\text{key}$ 
13          $y.\text{left} = z$ 
14     else  $y.\text{right} = z$ 
```

$\Theta(1)$

It visits a simple path from the root
to a leaf node. Hence $\Theta(h)$

$\Theta(1)$

Deletion

The procedure for deleting takes a BST T and a node z in T , and updates the pointers so that z does no longer belong to T while preserving the BST property.

Basic Idea:

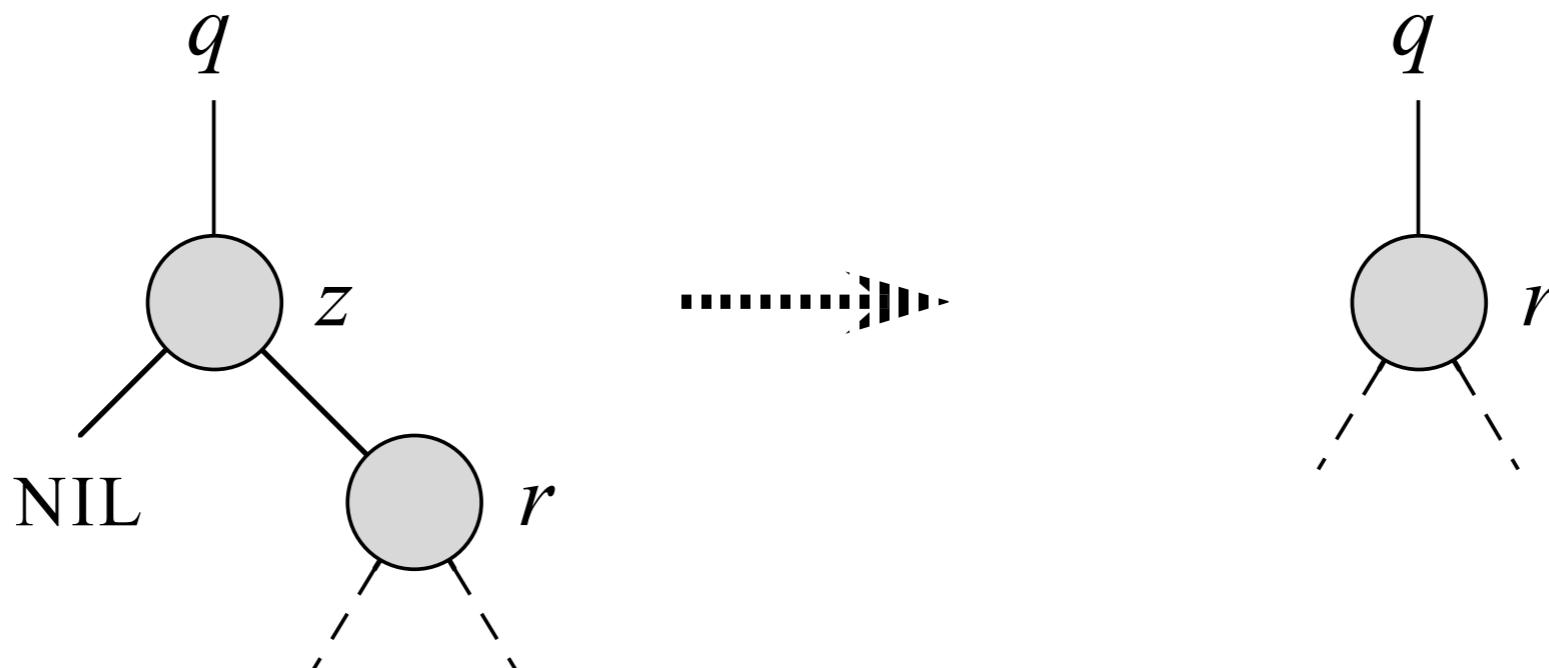
We consider three cases:

1. If z has no children we simply remove it
2. If z has just one child, then we elevate that child to take z 's position in the tree.
3. If z has both children, then we find z 's successor —which is in z 's right subtree and has no left subtree— and have it take z 's position in the tree

Tree-Delete

For the actual procedure, it will be more convenient to consider four cases instead:

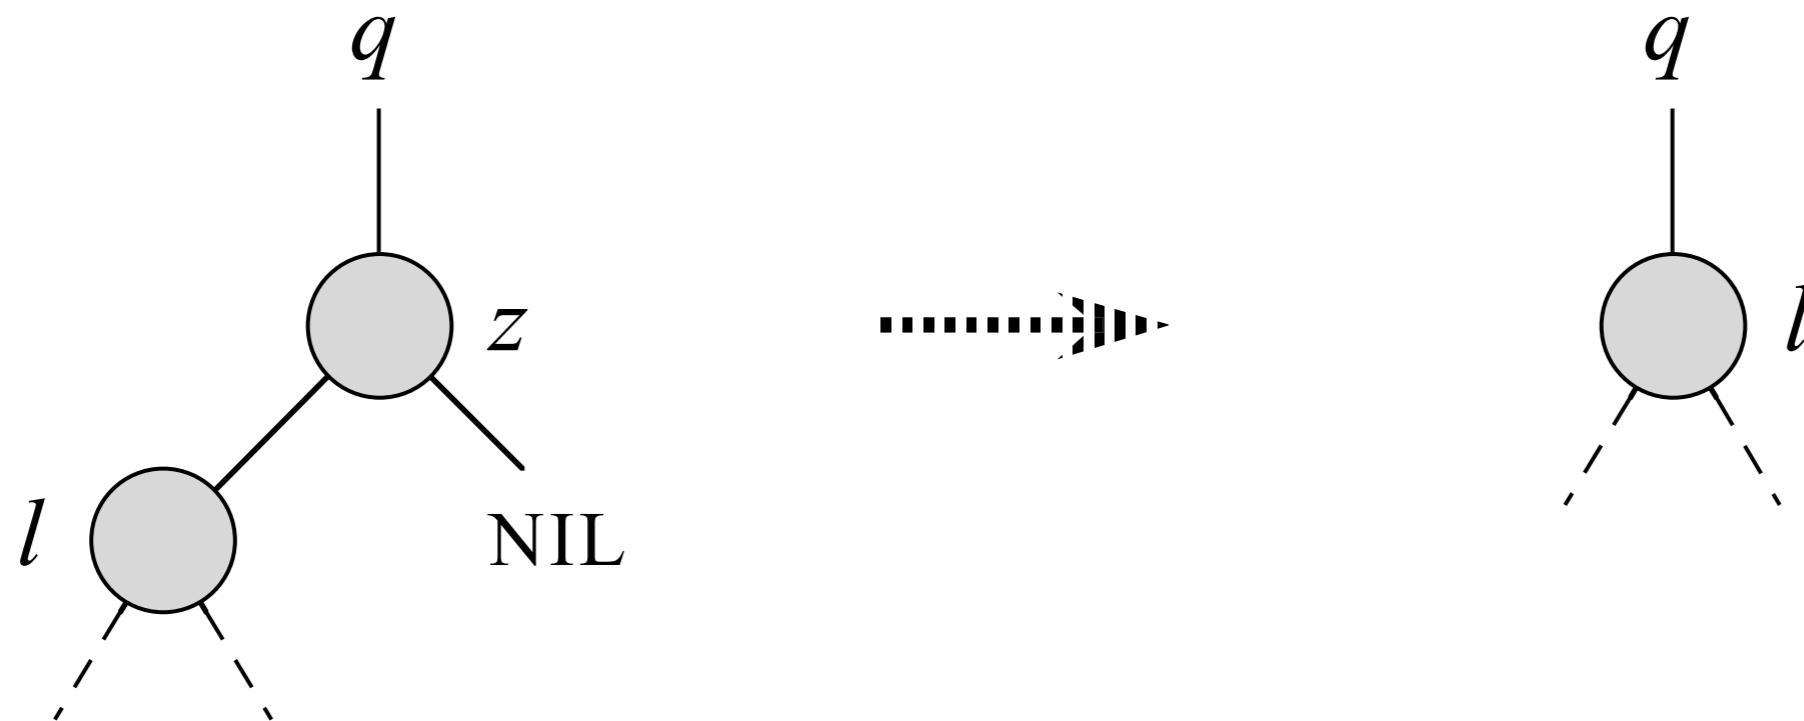
Case A: z has no left child. Then we replace z by its right child (which may be Nil)



Tree-Delete

For the actual procedure, it will be more convenient to consider four cases instead:

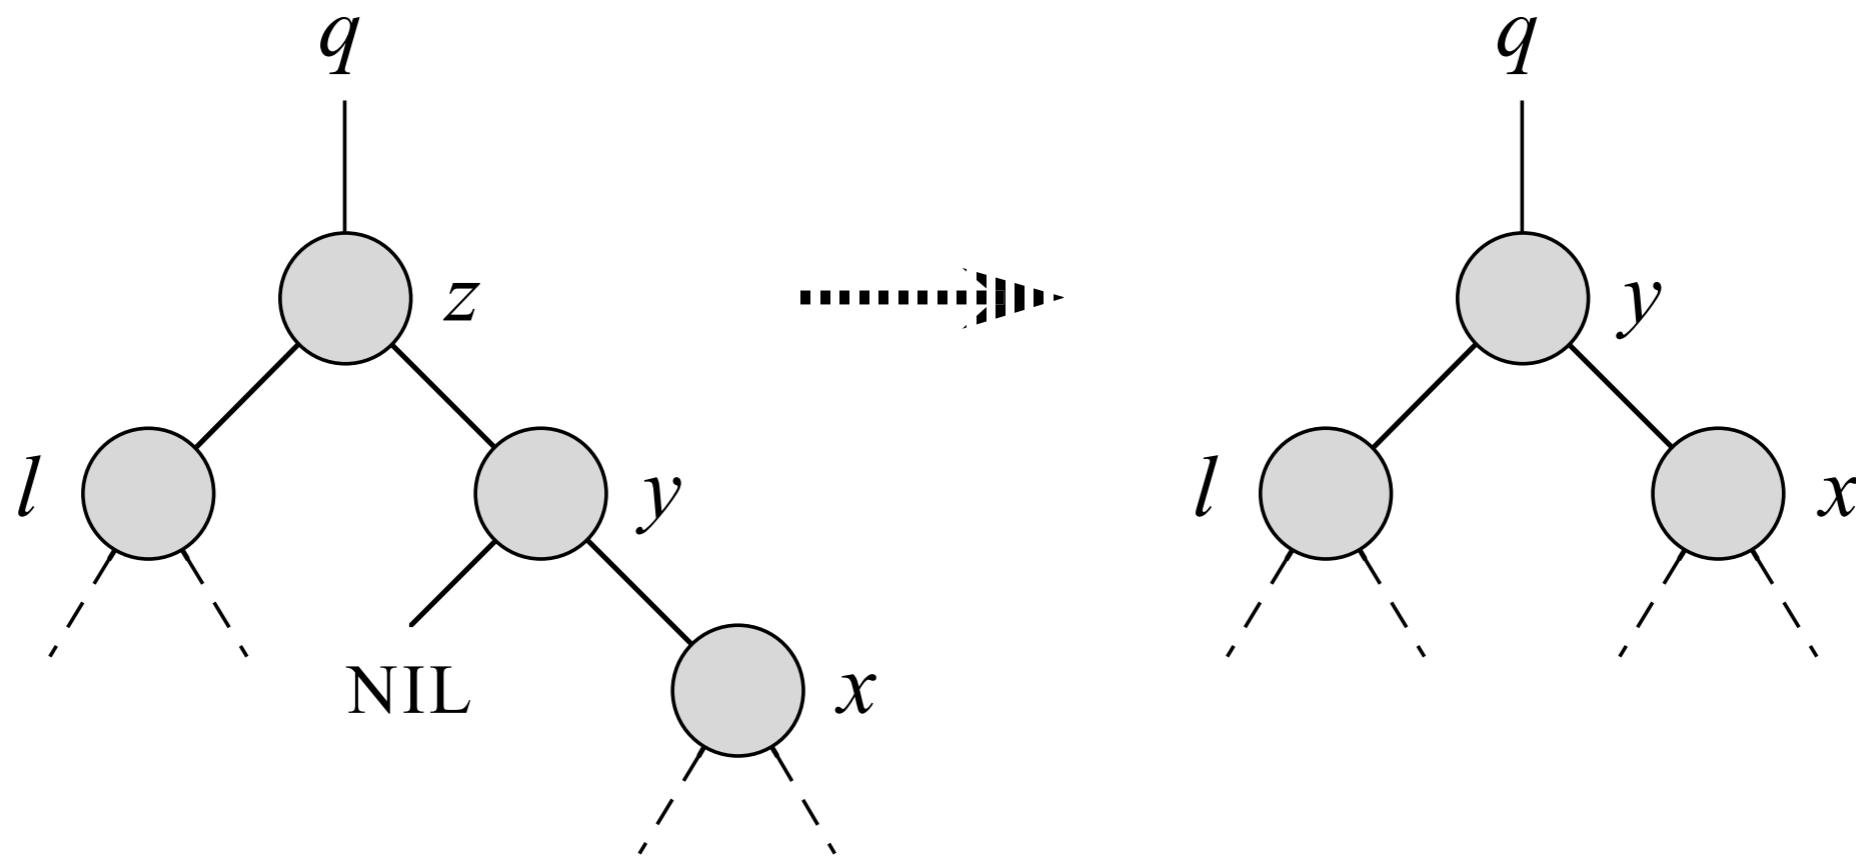
Case B: z has left child but no right child. Then we replace z by its left child



Tree-Delete

For the actual procedure, it will be more convenient to consider four cases instead:

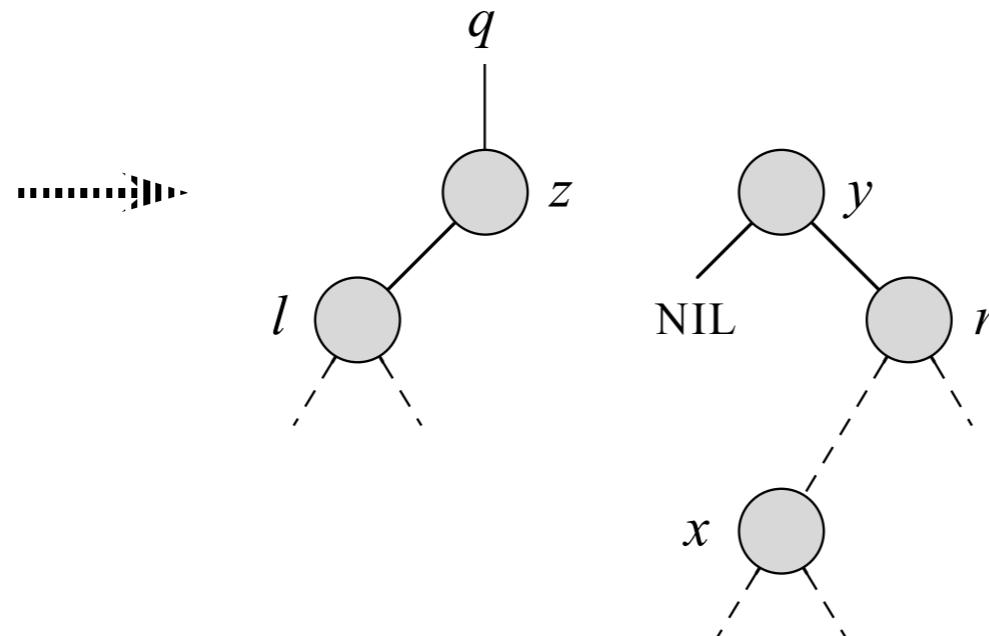
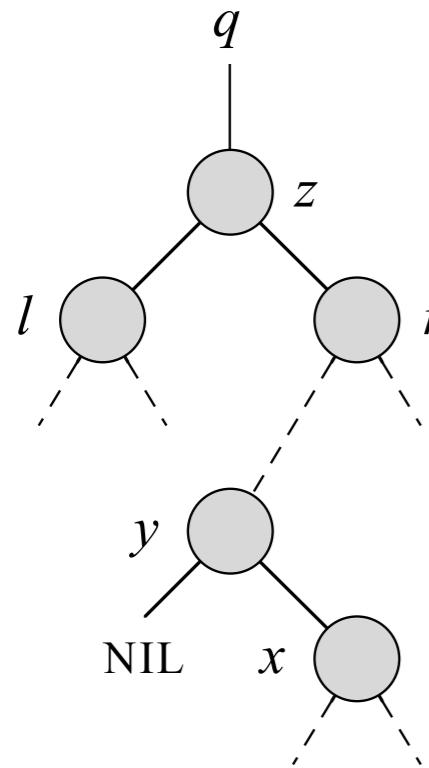
Case C: z has both children and z 's successor is its right child. Then “push” y up replacing z 's position



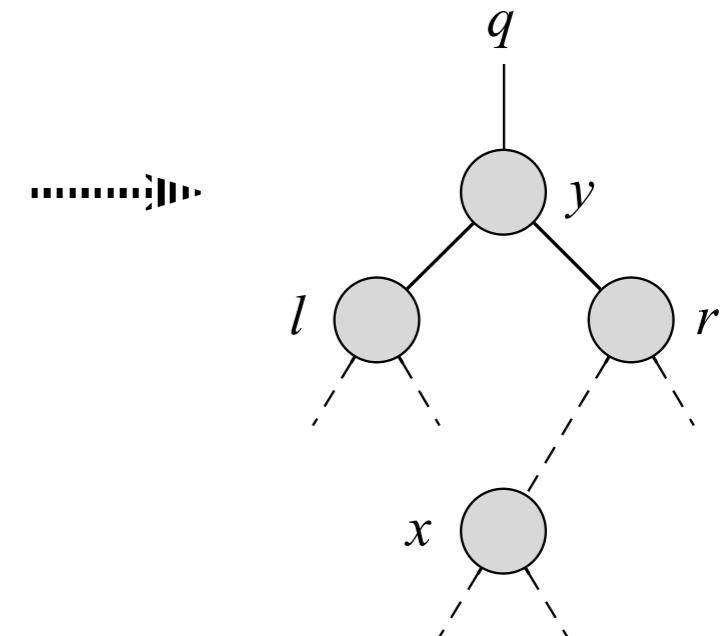
Tree-Delete

For the actual procedure, it will be more convenient to consider four cases instead:

Case D: z has both children and z 's successor is not its right child. Then, we first remove y from T , after we replace y with z 's position



D.1



D.2

The Transplant subroutine

- The pseudocode for **TREE-DELETE** will use a subroutine, called **TRANSPLANT**, that **replaces one subtree rooted at node u with the subtree rooted at node v .**
- Its worst-case **running-time** is $\Theta(1)$

TRANSPLANT(T, u, v)

```
1  if  $u.p == \text{NIL}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6  if  $v \neq \text{NIL}$ 
7       $v.p = u.p$ 
```

Node u 's parent ends up having
 v as its appropriate child

Node u 's parent becomes
node v 's parent

Tree-Delete: pseudocode

TREE-DELETE(T, z)

```
1  if  $z.left == \text{NIL}$ 
2      TRANSPLANT( $T, z, z.right$ )
3  elseif  $z.right == \text{NIL}$ 
4      TRANSPLANT( $T, z, z.left$ )
5  else  $y = \text{TREE-MINIMUM}(z.right)$ 
6      if  $y.p \neq z$ 
7          TRANSPLANT( $T, y, y.right$ )
8           $y.right = z.right$ 
9           $y.right.p = y$ 
10     TRANSPLANT( $T, z, y$ )
11      $y.left = z.left$ 
12      $y.left.p = y$ 
```

Case A

Case B

Case D.1

Case C

+

Case D.2

Tree-Delete: run-time

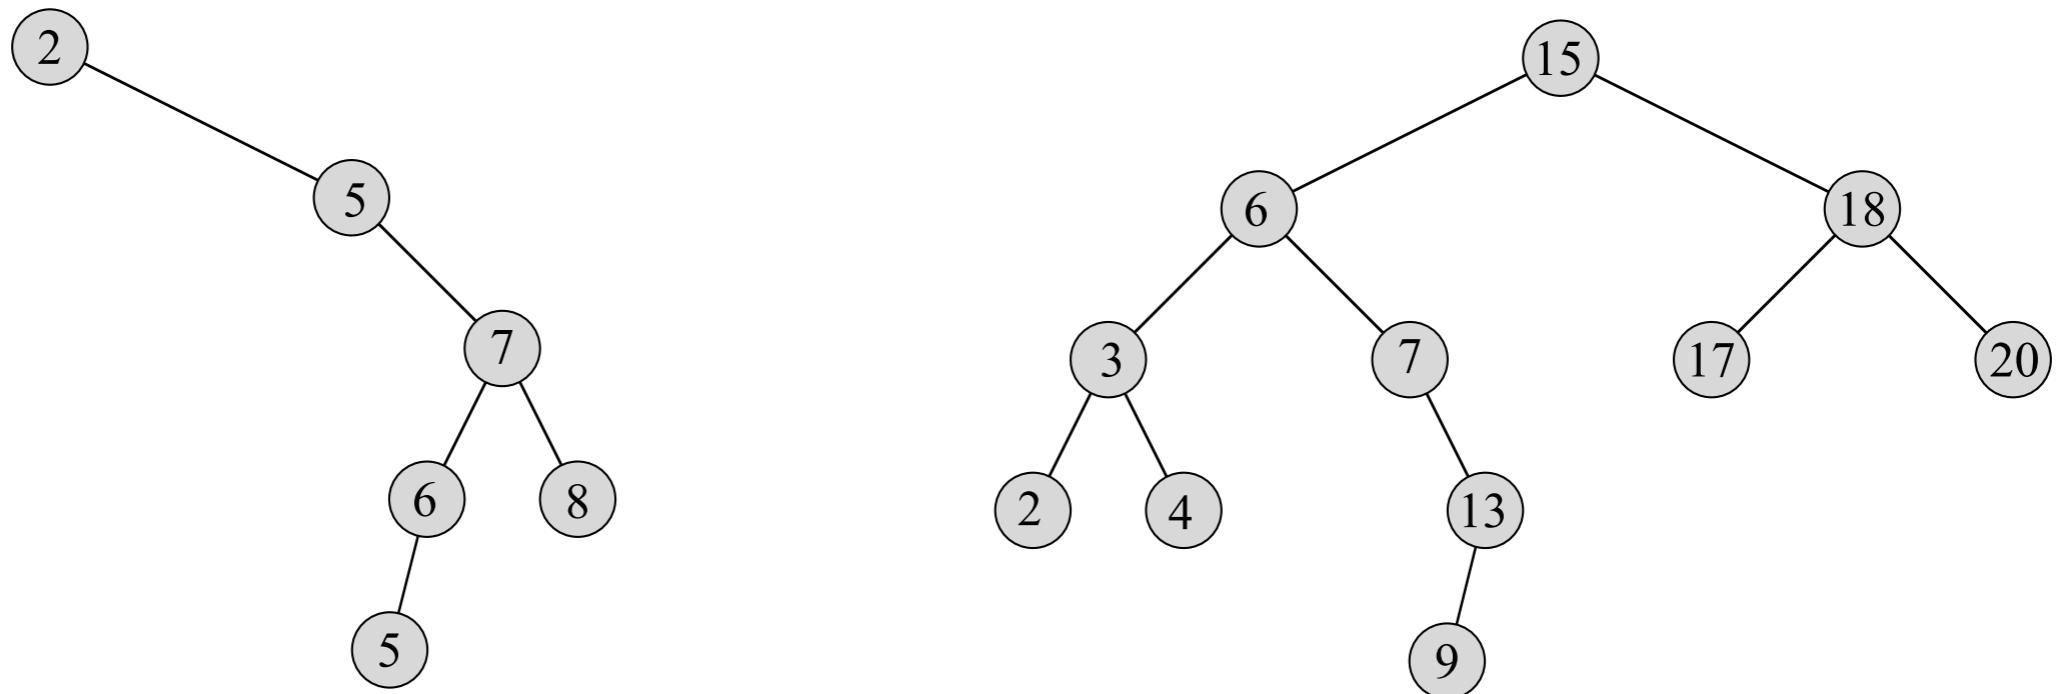
- **Worst-case** occurs when we are in Case C and Case D.
- Case A and Case B only take $\Theta(1)$

TREE-DELETE(T, z)

```
1  if  $z.left == \text{NIL}$ 
2      TRANSPLANT( $T, z, z.right$ )
3  elseif  $z.right == \text{NIL}$ 
4      TRANSPLANT( $T, z, z.left$ )
5  else  $y = \text{TREE-MINIMUM}(z.right)$  Θ( $h$ )
6      if  $y.p \neq z$ 
7          TRANSPLANT( $T, y, y.right$ )
8           $y.right = z.right$ 
9           $y.right.p = y$ 
10     TRANSPLANT( $T, z, y$ )
11      $y.left = z.left$ 
12      $y.left.p = y$ 
```

BST: Summary & Reflections

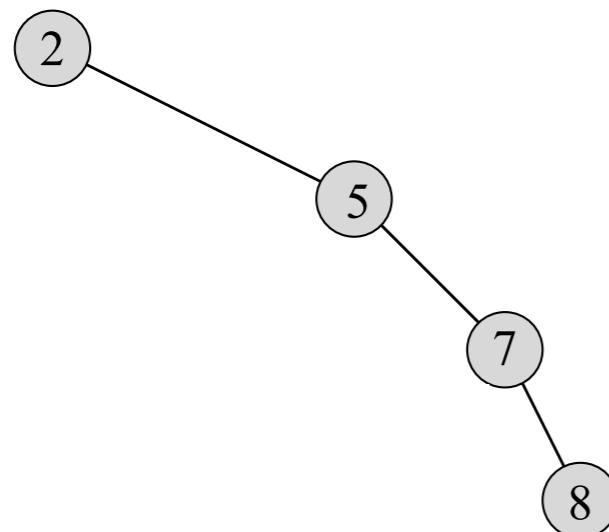
- BSTs are an effective data structure for dynamic sets
- **All the operations** we have seen **run in at most $\Theta(h)$ time**, where h is the height of the BST
 - **Good** when the tree is **balanced**: $h = \Theta(\lg n)$
 - **Bad** when the tree is **unbalanced**: $h = \Theta(n)$



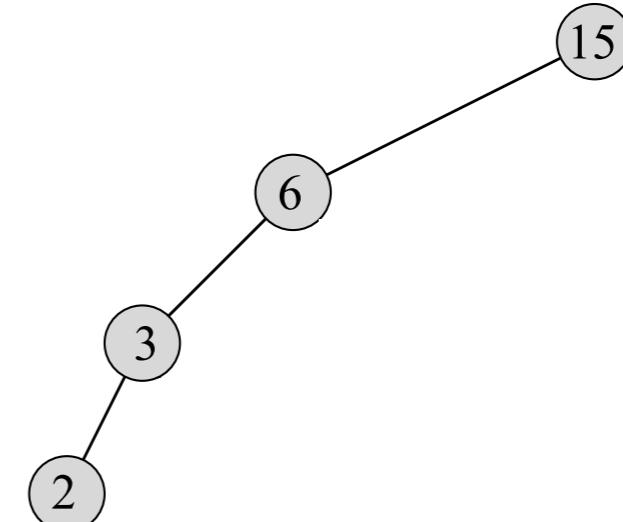
Balanced BSTs

- How easy is it to create an unbalanced tree by inserting a sequence of nodes in the tree?
 - If the sequence has increasing keys (reps. decreasing keys) the resulting tree boils down to a chain

Result of Insertion 2; 5; 7; 8



Result of Insertion 15; 6; 3; 2



Balanced BSTs

- Little is known about average height of a binary search tree when both insertion and deletion are used to create it.
- We can say something when the tree is created by insertion alone

Randomly-built BST:

n -keys BST obtained by *inserting the keys in random order* into an initially empty tree, where *each of the $n!$ permutations of the input keys are equally likely*

Theorem 12.4

The expected height of a randomly built binary search tree on n distinct keys is $O(\lg n)$.

Red-Black Trees

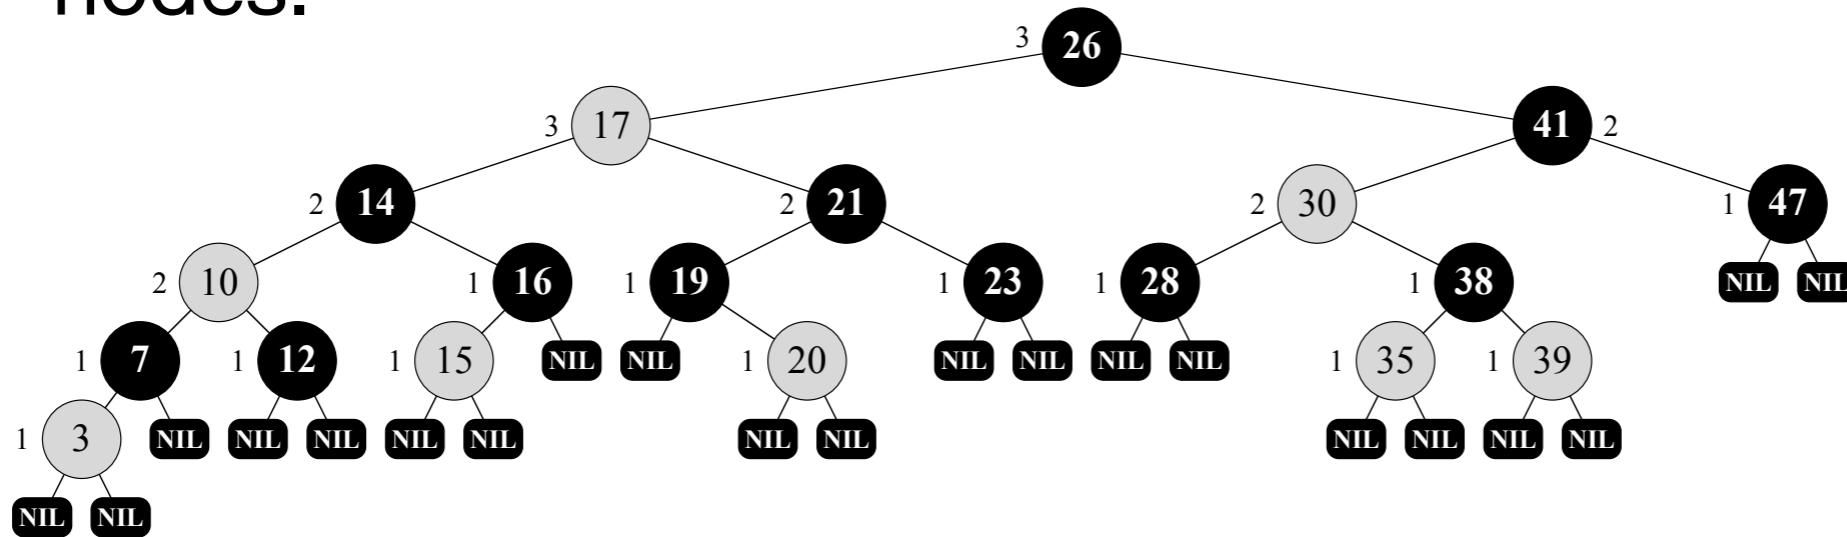
Red-Black Trees

- Red-back Trees **are BSTs** where each node has an additional colour attribute which can be either Red or Black
- By constraining node colours on any simple path from the root to a leaf, red-black trees ensure that no such path is more than twice as long as any other
- Red-black trees are approximately balanced, that is $h = O(\lg n)$

Red-Black Trees

A red-black tree is a BST that additionally satisfies the following **red-black properties**:

1. Every node is either red or black
2. The root is black
3. Every leaf (**NIL**) is black
4. If a node is red, then both its children are black
5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.



Red-Black Trees

Lemma 13.1

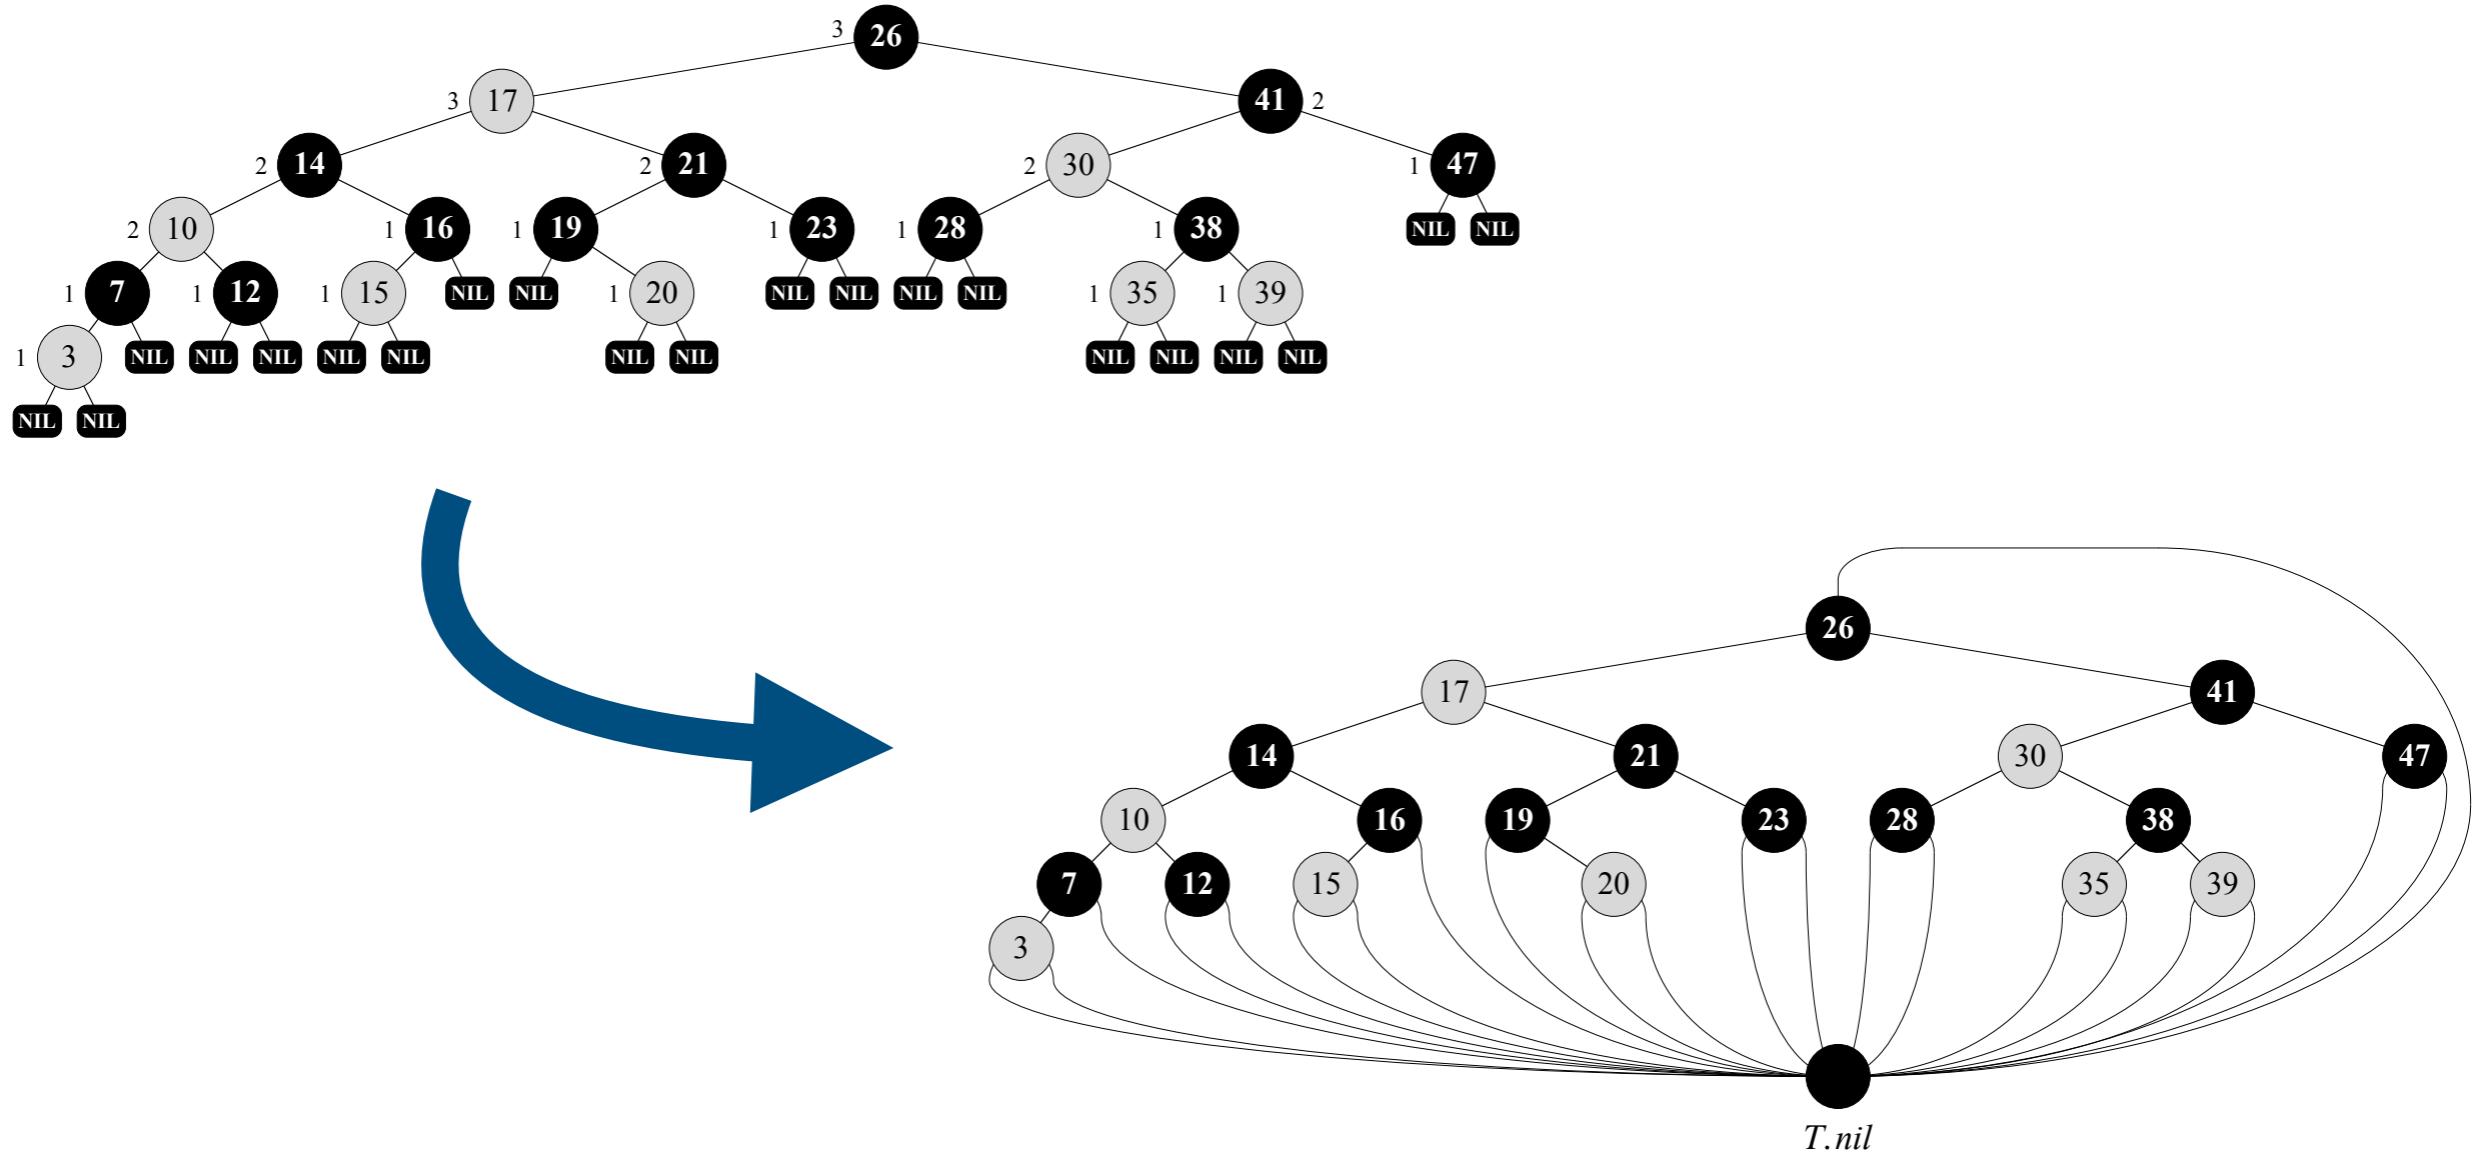
A red-black tree with n internal nodes has height at most $2 \lg(n + 1)$.

Optional: read the proof of Lemma 13.1 in CLRS

- Remember that red-black trees are special BSTs.
- Thus Lemma 13.1 implies that we can implement the dynamic-set operations **SEARCH**, **MINIMUM**, **MAXIMUM**, **SUCCESSOR**, and **PREDECESSOR** in $O(\lg n)$ time, where n is the number of elements in the tree
- Insertion and deletion require more work to preserve the red-black-tree property

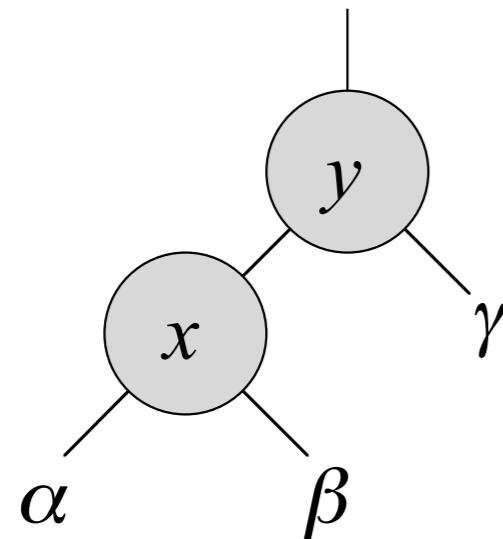
Red-Black Trees

- It will be convenient to replace Nil pointers with *nil* nodes
- In fact, we only need one such node which can be shared



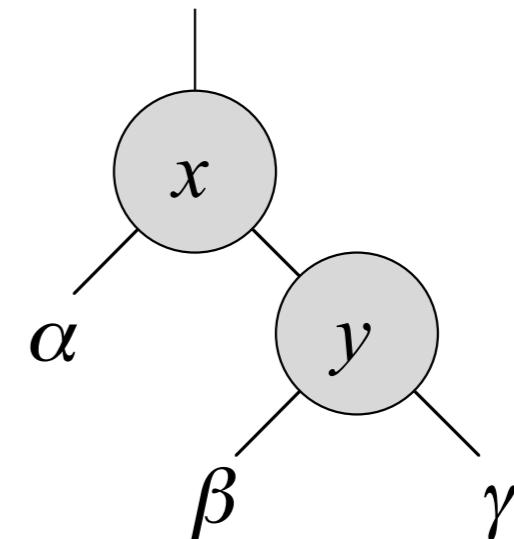
Rotations

- Insertion and deletion in red-black trees have to preserve the red-black tree property
- This is done by
 - Changing the colours of some nodes in the tree, and
 - Changing the pointer structure through rotations
- Rotation is a local operation that **preserves the BST property** (may not preserve the red-black tree property)



LEFT-ROTATE(T, x)

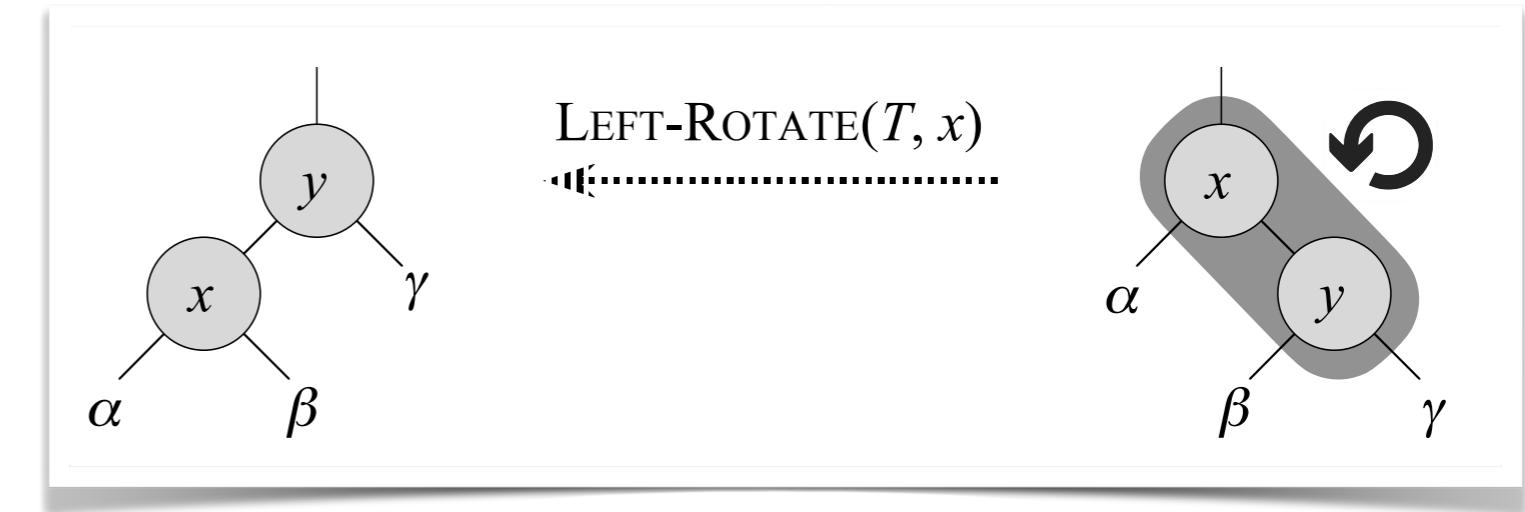
RIGHT-ROTATE(T, y)



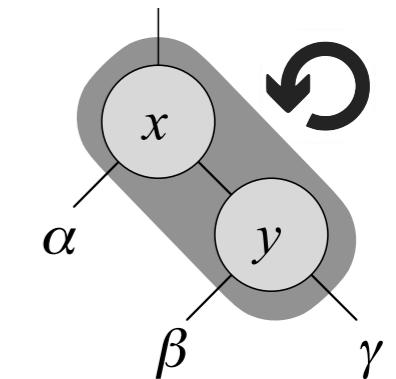
Left-Rotate: pseudocode

LEFT-ROTATE(T, x)

```
1   $y = x.right$ 
2   $x.right = y.left$ 
3  if  $y.left \neq T.nil$ 
4     $y.left.p = x$ 
5
6   $y.p = x.p$ 
7  if  $x.p == T.nil$ 
8     $T.root = y$ 
9  elseif  $x == x.p.left$ 
10    $x.p.left = y$ 
11  else  $x.p.right = y$ 
12   $y.left = x$ 
13   $x.p = y$ 
```



LEFT-ROTATE(T, x)



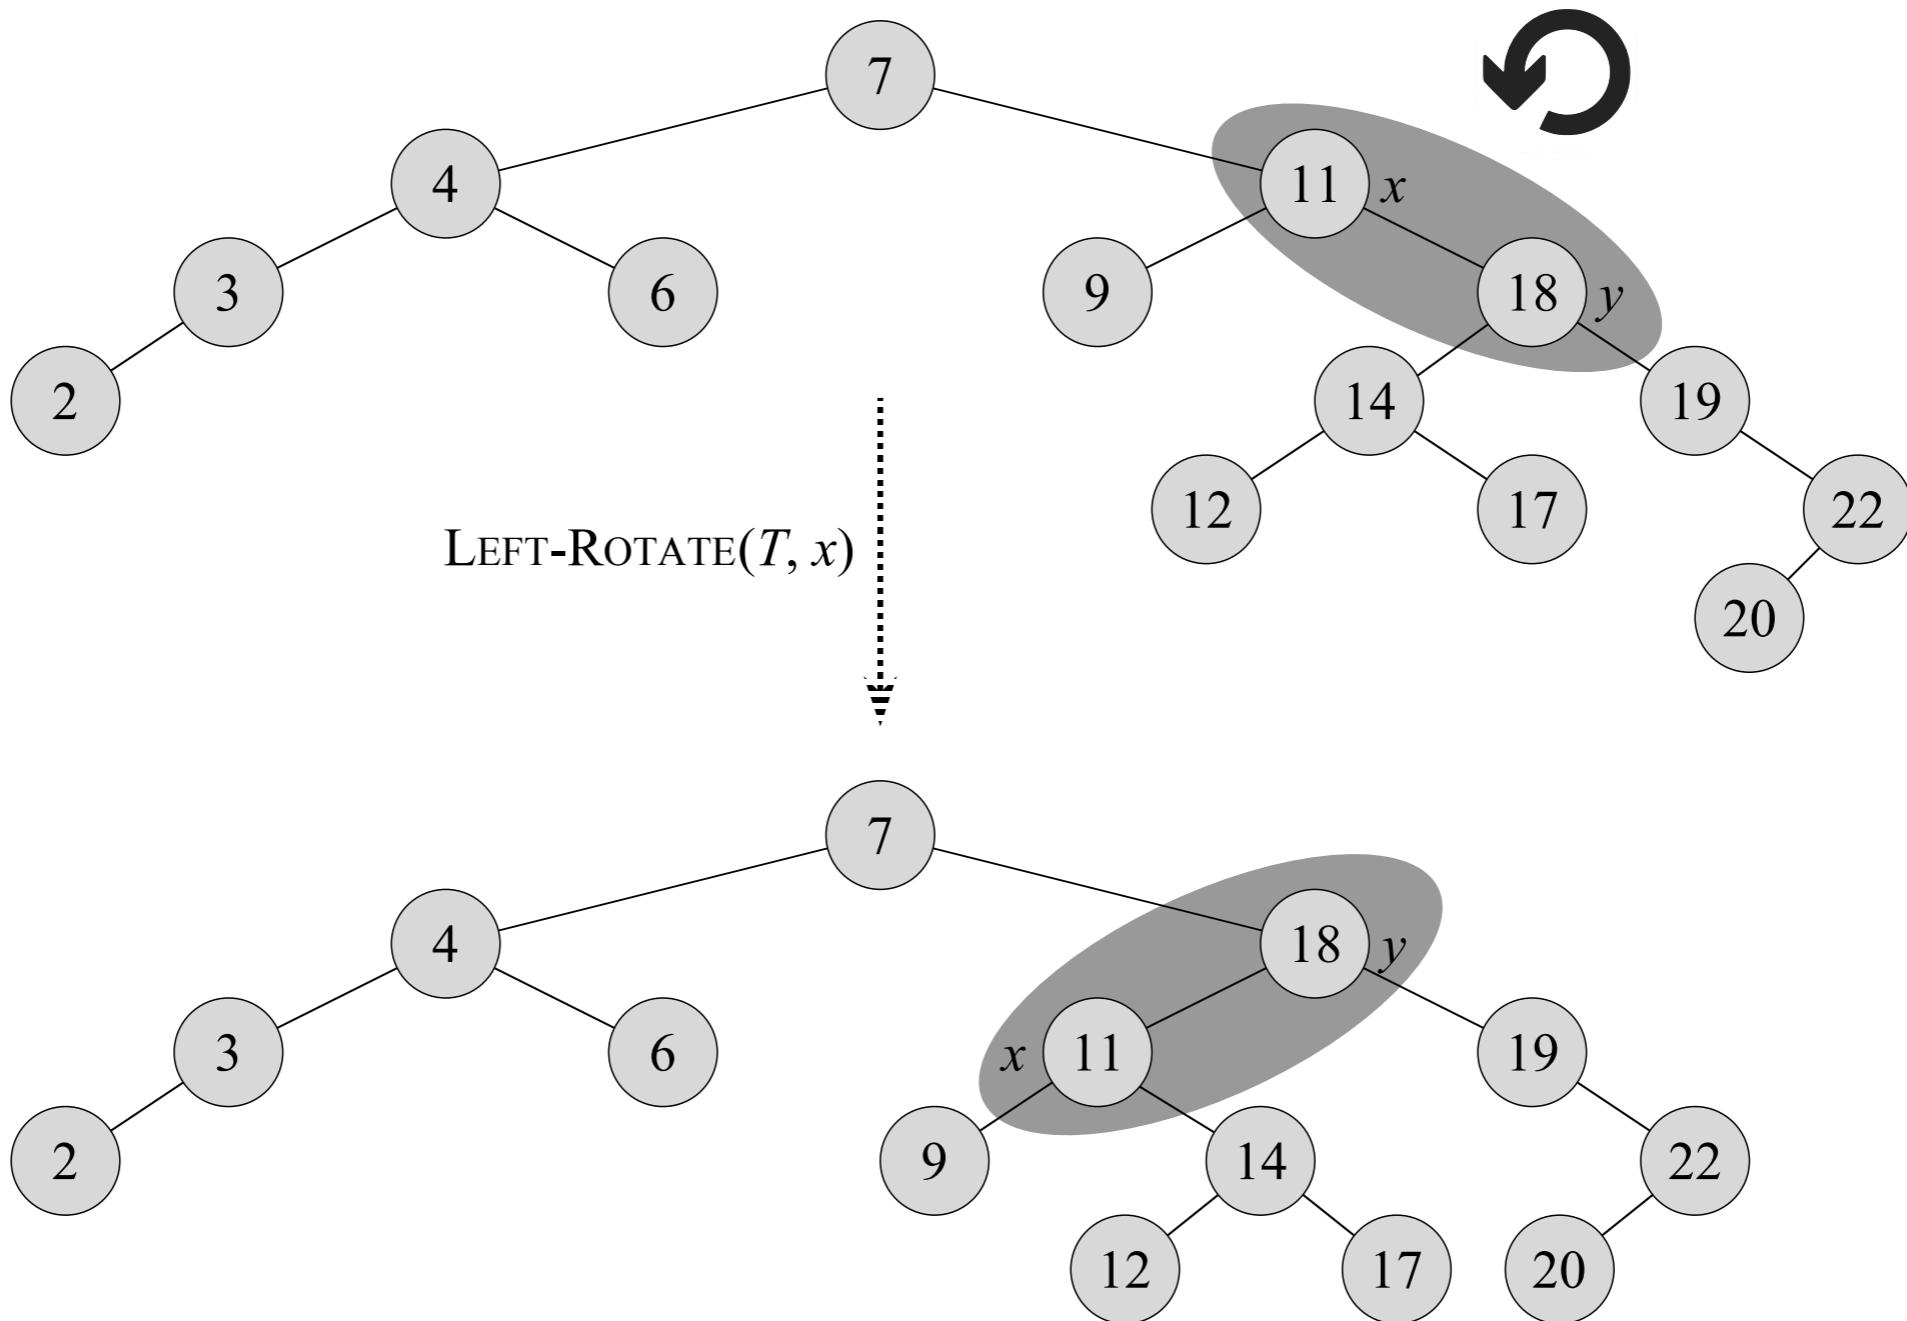
set y

Turn y 's left subtree (i.e., β)
into x 's right subtree

Link x 's parent to y

Put x on y 's left

Left-Rotate: example



RB-Insert

RB-INSERT(T, z)

```
1   $y = T.nil$ 
2   $x = T.root$ 
3  while  $x \neq T.nil$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8       $z.p = y$ 
9      if  $y == T.nil$ 
10          $T.root = z$ 
11     elseif  $z.key < y.key$ 
12          $y.left = z$ 
13     else  $y.right = z$ 
14      $z.left = T.nil$ 
15      $z.right = T.nil$ 
16      $z.color = RED$ 
17     RB-INSERT-FIXUP( $T, z$ )
```

TREE-INSERT(T, z)

```
1   $y = NIL$ 
2   $x = T.root$ 
3  while  $x \neq NIL$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8       $z.p = y$ 
9      if  $y == NIL$ 
10          $T.root = z$ 
11     elseif  $z.key < y.key$ 
12          $y.left = z$ 
13     else  $y.right = z$ 
```

Set z as **red** and set its children

Restore the RB property!

How to restore RB property?

To understand how RB-Insert-Fixup works we shall first determine what violations of the RB property are introduced in RB-Insert when z is inserted an coloured red

Red-black properties:

1. Every node is either red or black 
2. The root is black 
3. Every leaf (`NIL`) is black 
4. If a node is red, then both its children are black 
5. For each node, all simple paths from the node to descendant leaves contains the same number of black nodes. 

How to restore RB property?

Basic Idea

We start from the inserted element z (which is initially at the bottom of the tree) and we move up the violations of the RB property up to the root

Loop Invariant:

1. Node z is **red**
2. If z 's parent (i.e., $z.p$) is the root, then $z.p$ is black
3. If the tree violates any of the RB properties, then
 - Either it violates Prop. 2 (because z is the root), or
 - It violates Prop. 4 (because also z 's parent is red)

RB-Insert-Fixup

RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == \text{RED}$ 
2    if  $z.p == z.p.p.left$ 
3       $y = z.p.p.right$ 
4      if  $y.color == \text{RED}$ 
5         $z.p.color = \text{BLACK}$ 
6         $y.color = \text{BLACK}$ 
7         $z.p.p.color = \text{RED}$ 
8         $z = z.p.p$ 
9    else if  $z == z.p.right$ 
10       $z = z.p$ 
11      LEFT-ROTATE( $T, z$ )
12       $z.p.color = \text{BLACK}$ 
13       $z.p.p.color = \text{RED}$ 
14      RIGHT-ROTATE( $T, z.p.p$ )
15    else (same as then clause
16      with “right” and “left” exchanged)
17
18   $T.root.color = \text{BLACK}$ 
```

The while loop moves the violation up in the tree preserving the invariant

RB-Insert-Fixup

RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == \text{RED}$ 
2    if  $z.p == z.p.p.left$ 
3       $y = z.p.p.right$ 
4      if  $y.color == \text{RED}$ 
5         $z.p.color = \text{BLACK}$ 
6         $y.color = \text{BLACK}$ 
7         $z.p.p.color = \text{RED}$ 
8         $z = z.p.p$ 
9      else if  $z == z.p.right$ 
10         $z = z.p$ 
11        LEFT-ROTATE( $T, z$ )
12         $z.p.color = \text{BLACK}$ 
13         $z.p.p.color = \text{RED}$ 
14        RIGHT-ROTATE( $T, z.p.p$ )
15    else (same as then clause
16      with “right” and “left” exchanged)
17
18   $T.root.color = \text{BLACK}$ 
```

The while loop moves the violation up in the tree preserving the invariant

There are two symmetric procedures respectively if z 's parent is a left child or a right child

RB-Insert-Fixup

RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == \text{RED}$ 
2    if  $z.p == z.p.p.left$ 
3       $y = z.p.p.right$ 
4      if  $y.color == \text{RED}$ 
5         $z.p.color = \text{BLACK}$ 
6         $y.color = \text{BLACK}$ 
7         $z.p.p.color = \text{RED}$ 
8         $z = z.p.p$ 
9      else if  $z == z.p.right$ 
10         $z = z.p$ 
11        LEFT-ROTATE( $T, z$ )
12         $z.p.color = \text{BLACK}$ 
13         $z.p.p.color = \text{RED}$ 
14        RIGHT-ROTATE( $T, z.p.p$ )
15      else (same as then clause
16          with "right" and "left" exchanged)
17
18       $T.root.color = \text{BLACK}$ 
```

The while loop moves the violation up in the tree preserving the invariant

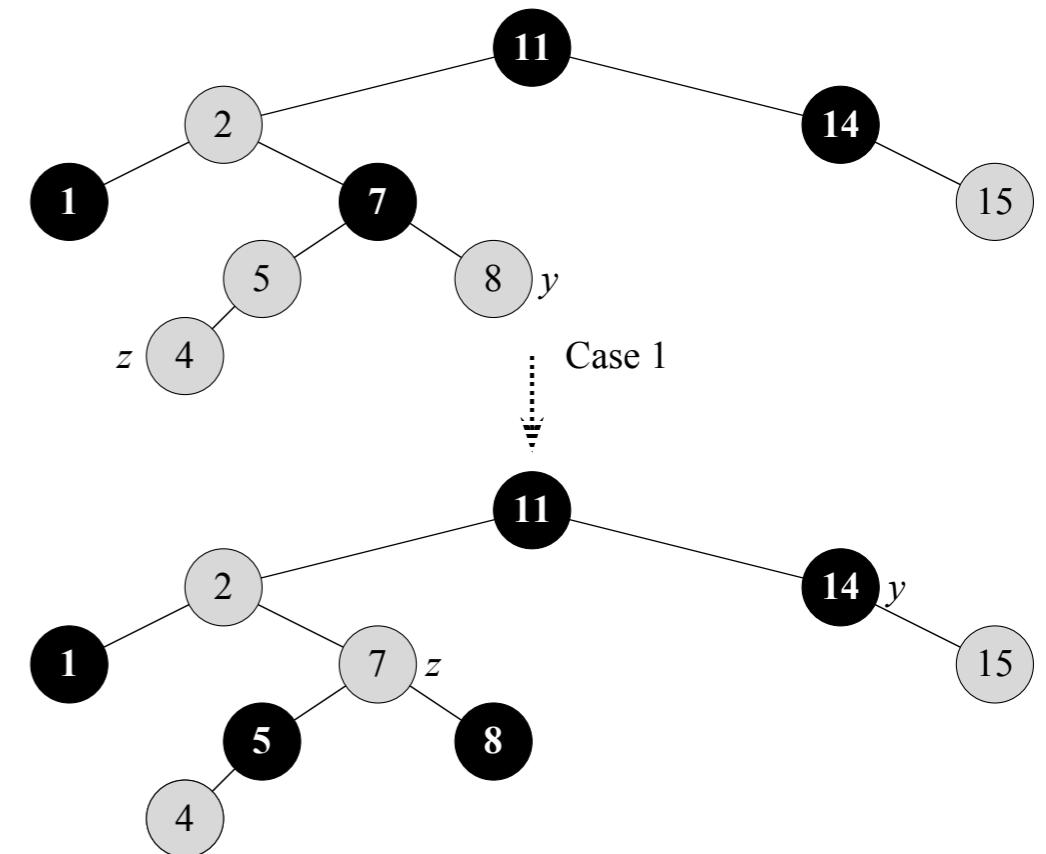
There are two symmetric procedures respectively if z 's parent is a left child or a right child

Each are subdivided in 3 cases (where in particular Case 2 redirects to Case 3)

RB-Insert-Fixup: Case 1

RB-INSERT-FIXUP(T, z)

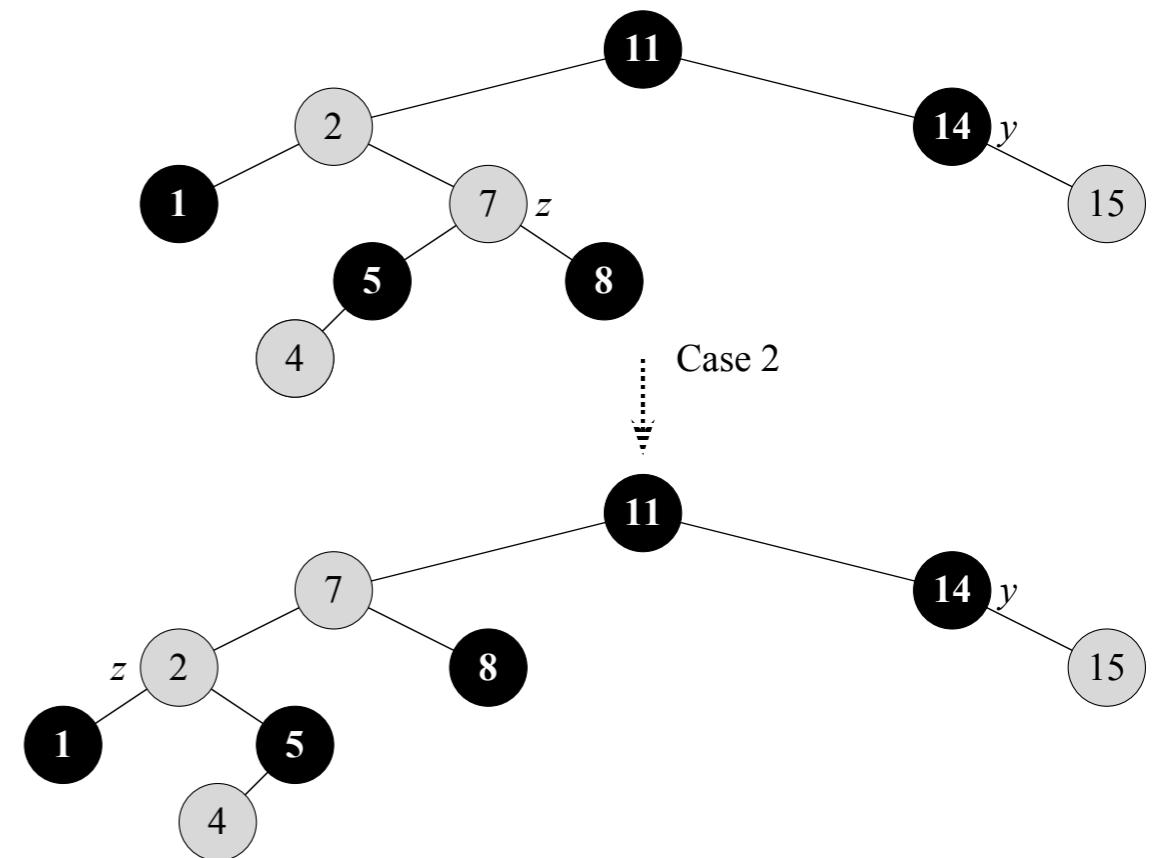
```
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9      else if  $z == z.p.right$ 
10          $z = z.p$ 
11         LEFT-ROTATE( $T, z$ )
12          $z.p.color = \text{BLACK}$ 
13          $z.p.p.color = \text{RED}$ 
14         RIGHT-ROTATE( $T, z.p.p$ )
15     else (same as then clause
           with “right” and “left” exchanged)
16      $T.root.color = \text{BLACK}$ 
```



RB-Insert-Fixup: Case 2

RB-INSERT-FIXUP(T, z)

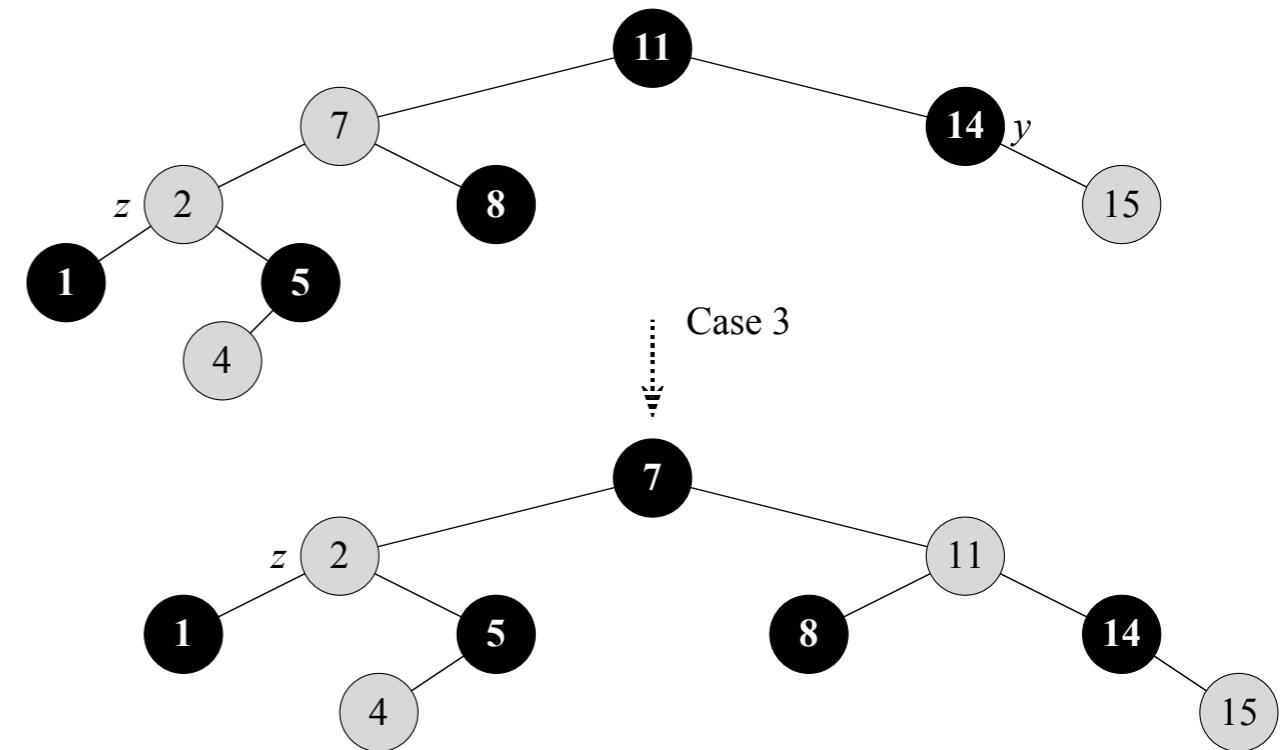
```
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9      else if  $z == z.p.right$ 
10          $z = z.p$ 
11         LEFT-ROTATE( $T, z$ )
12          $z.p.color = \text{BLACK}$ 
13          $z.p.p.color = \text{RED}$ 
14         RIGHT-ROTATE( $T, z.p.p$ )
15     else (same as then clause
           with “right” and “left” exchanged)
16      $T.root.color = \text{BLACK}$ 
```



RB-Insert-Fixup: Case 3

RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9      else if  $z == z.p.right$ 
10          $z = z.p$ 
11         LEFT-ROTATE( $T, z$ )
12          $z.p.color = \text{BLACK}$ 
13          $z.p.p.color = \text{RED}$ 
14         RIGHT-ROTATE( $T, z.p.p$ )
15     else (same as then clause
           with “right” and “left” exchanged)
16      $T.root.color = \text{BLACK}$ 
```



RB-Insert: run-time

Worst-case running time $O(\lg n)$ because $h = O(\lg n)$

RB-INSERT(T, z)

```
1   $y = T.nil$ 
2   $x = T.root$ 
3  while  $x \neq T.nil$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$ 
9  if  $y == T.nil$ 
10      $T.root = z$ 
11  elseif  $z.key < y.key$ 
12       $y.left = z$ 
13  else  $y.right = z$ 
14   $z.left = T.nil$ 
15   $z.right = T.nil$ 
16   $z.color = \text{RED}$ 
17  RB-INSERT-FIXUP( $T, z$ )
```

$\Theta(h)$ as insertion in BST

$\Theta(1)$

$\Theta(h)$. The number of loops is bounded by h

RB-Delete

RB-DELETE(T, z)

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4     $x = z.\text{right}$ 
5    RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7     $x = z.\text{left}$ 
8    RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10    $y\text{-original-color} = y.\text{color}$ 
11    $x = y.\text{right}$ 
12   if  $y.p == z$ 
13      $x.p = y$ 
14   else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15      $y.\text{right} = z.\text{right}$ 
16      $y.\text{right}.p = y$ 
17   RB-TRANSPLANT( $T, z, y$ )
18    $y.\text{left} = z.\text{left}$ 
19    $y.\text{left}.p = y$ 
20    $y.\text{color} = z.\text{color}$ 
21 if  $y\text{-original-color} == \text{BLACK}$ 
22   RB-DELETE-FIXUP( $T, x$ )
```

TREE-DELETE(T, z)

```
1  if  $z.\text{left} == \text{NIL}$ 
2    TRANSPLANT( $T, z, z.\text{right}$ )
3  elseif  $z.\text{right} == \text{NIL}$ 
4    TRANSPLANT( $T, z, z.\text{left}$ )
5  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
6    if  $y.p \neq z$ 
7      TRANSPLANT( $T, y, y.\text{right}$ )
8       $y.\text{right} = z.\text{right}$ 
9       $y.\text{right}.p = y$ 
10     TRANSPLANT( $T, z, y$ )
11      $y.\text{left} = z.\text{left}$ 
12      $y.\text{left}.p = y$ 
```

+

Some additional code which
keeps track of two nodes y and x

Restore the RB
property!

RB-Delete

RB-DELETE(T, z)

```

1    $y = z$ 
2    $y\text{-original-color} = y.\text{color}$ 
3   if  $z.\text{left} == T.\text{nil}$ 
4      $x = z.\text{right}$ 
5     RB-TRANSPLANT( $T, z, z.\text{right}$ )
6   elseif  $z.\text{right} == T.\text{nil}$ 
7      $x = z.\text{left}$ 
8     RB-TRANSPLANT( $T, z, z.\text{left}$ )
9   else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10   $y\text{-original-color} = y.\text{color}$ 
11   $x = y.\text{right}$ 
12  if  $y.p == z$ 
13     $x.p = y$ 
14  else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15     $y.\text{right} = z.\text{right}$ 
16     $y.\text{right}.p = y$ 
17  RB-TRANSPLANT( $T, z, y$ )
18   $y.\text{left} = z.\text{left}$ 
19   $y.\text{left}.p = y$ 
20   $y.\text{color} = z.\text{color}$ 
21  if  $y\text{-original-color} == \text{BLACK}$ 
22    RB-DELETE-FIXUP( $T, x$ )

```

If z 's left (or right) child is Nil

- $y = z$ is the node removed, and
- x is the node that takes y 's old position in T

If z has both children Nil

- y is the node that takes z 's position, and
- x is the node that takes y 's old position in T

- Here y has been removed or has taken z 's place and colour
- If y 's original colour was black then some RB property may be violated

How to restore RB property?

To understand how RB-Delete-Fixup works we shall first determine what violations of the RB property are introduced in RB-Delete when y is deleted (i.e., $y = z$) or moved in z 's original position.

If y 's original colour was **red**, then no RB property is violated

Red-black properties:

1. Every node is either red or black ✓
2. The root is black ✓
3. Every leaf (**NIL**) is black ✓
4. If a node is red, then both its children are black ✓
5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes. ✓

y was not the root

- y takes z 's place in T , along with z 's colour, we cannot have two adjacent red nodes at y 's new position in the tree.
- In addition, if y was not z 's right child, then y 's original right child x replaces y in T . Since y was red, then x must be black.

No black heights have changed

How to restore RB property?

To understand how RB-Delete-Fixup works we shall first determine what violations of the RB property are introduced in RB-Delete when y is deleted (i.e., $y = z$) or moved in z 's original position.

If y 's original colour was **Black**

Red-black properties:

1. Every node is either red or black
2. The root is black
3. Every leaf (`NIL`) is black
4. If a node is red, then both its children are black
5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.

If y was the root, x may be red and take its position

If x and x 's parent are both red after x has replaced y 's original position

Moving y makes any simple path previously contain y to have one fewer black node

How to restore RB property?

To understand how RB-Delete-Fixup works we shall first determine what violations of the RB property are introduced in RB-Delete when y is deleted (i.e., $y = z$) or moved in z 's original position.

If y 's original colour was **Black**

Red-black properties:

1. Every node is either red or black ✖
2. The root is black ✖
3. Every leaf (`NIL`) is black ✓
4. If a node is red, then both its children are black ✖
5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes. ✓

We can fictionally ‘fix’ Prop. 5 by counting one more black token at x (who replaced y 's position). Hence, x can be **‘doubly black’** or **‘red-and-black’**.

RB-Delete-Fixup

Basic Idea

We start from x and we **move up the extra black token** up the tree until:

- x points to a red-and-black node, in which case we colour x just black
- x points to the root, in which case we colour x just black; or
- Having performed suitable rotations and recolouring, the loop exits

RB-Delete-Fixup

RB-DELETE-FIXUP(T, x)

```
1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22      else (same as then clause with "right" and "left" exchanged)
23       $x.color = \text{BLACK}$ 
```

The while loop moves the extra black token (pointed by x) up in the tree

RB-Delete-Fixup

RB-DELETE-FIXUP(T, x)

```
1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22    else (same as then clause with "right" and "left" exchanged)
23     $x.color = \text{BLACK}$ 
```

The while loop moves the extra black token (pointed by x) up in the tree

There are two symmetric procedures respectively if x is a left child or a right child

RB-Delete-Fixup

RB-DELETE-FIXUP(T, x)

```
1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22    else (same as then clause with "right" and "left" exchanged)
23     $x.color = \text{BLACK}$ 
```

Case 1

Case 2

Case 3

Case 4

The while loop moves the extra black token (pointed by x) up in the tree

There are two symmetric procedures respectively if x is a left child or a right child

Each divided in 4 cases

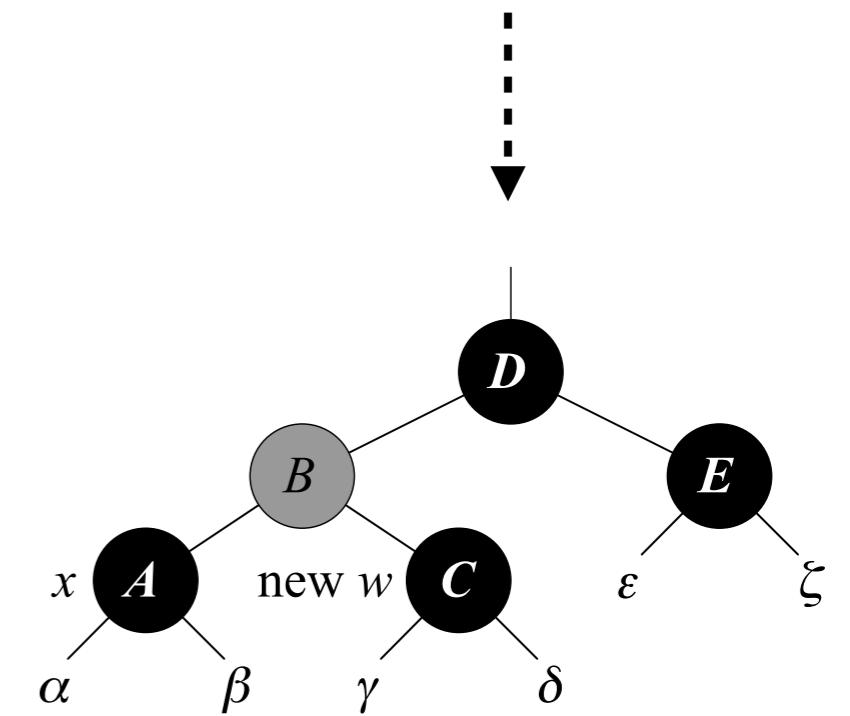
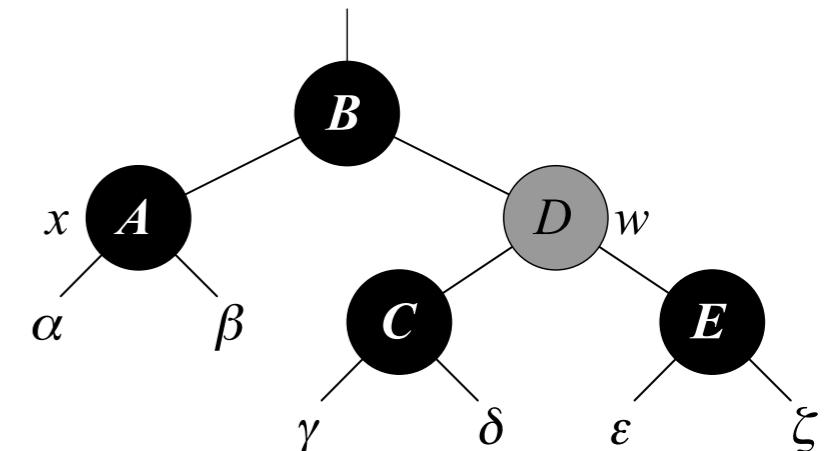
RB-Delete-Fixup: Case 1

RB-DELETE-FIXUP(T, x)

```

1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22    else (same as then clause with “right” and “left” exchanged)
23     $x.color = \text{BLACK}$ 
```

x 's sibling w is red



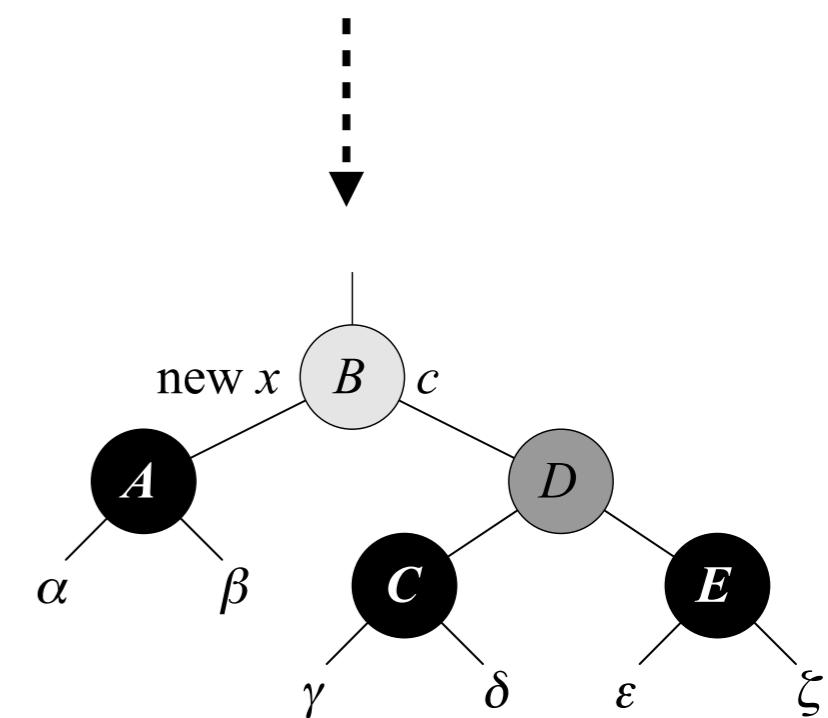
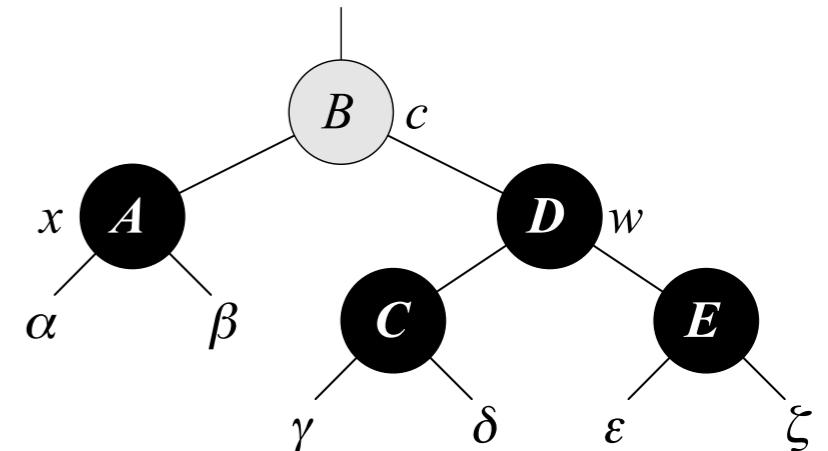
RB-Delete-Fixup: Case 2

RB-DELETE-FIXUP(T, x)

```

1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22      else (same as then clause with “right” and “left” exchanged)
23       $x.color = \text{BLACK}$ 
```

x 's sibling w is black, and both of w 's children are black



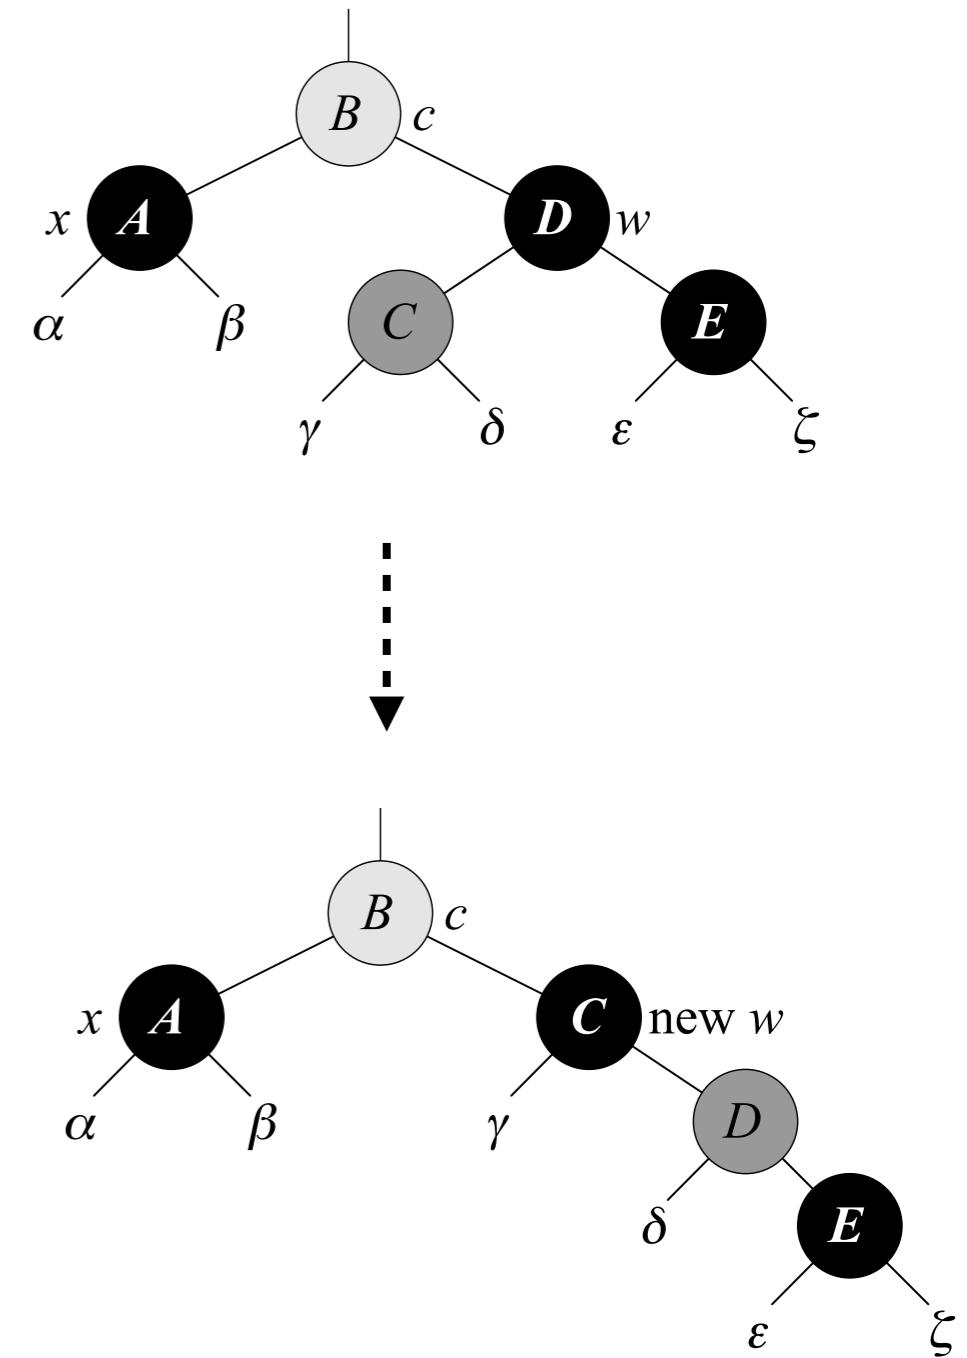
RB-Delete-Fixup: Case 3

RB-DELETE-FIXUP(T, x)

```

1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22      else (same as then clause with “right” and “left” exchanged)
23       $x.color = \text{BLACK}$ 
```

x 's sibling w is black, w 's left child is red, and w 's right child is black



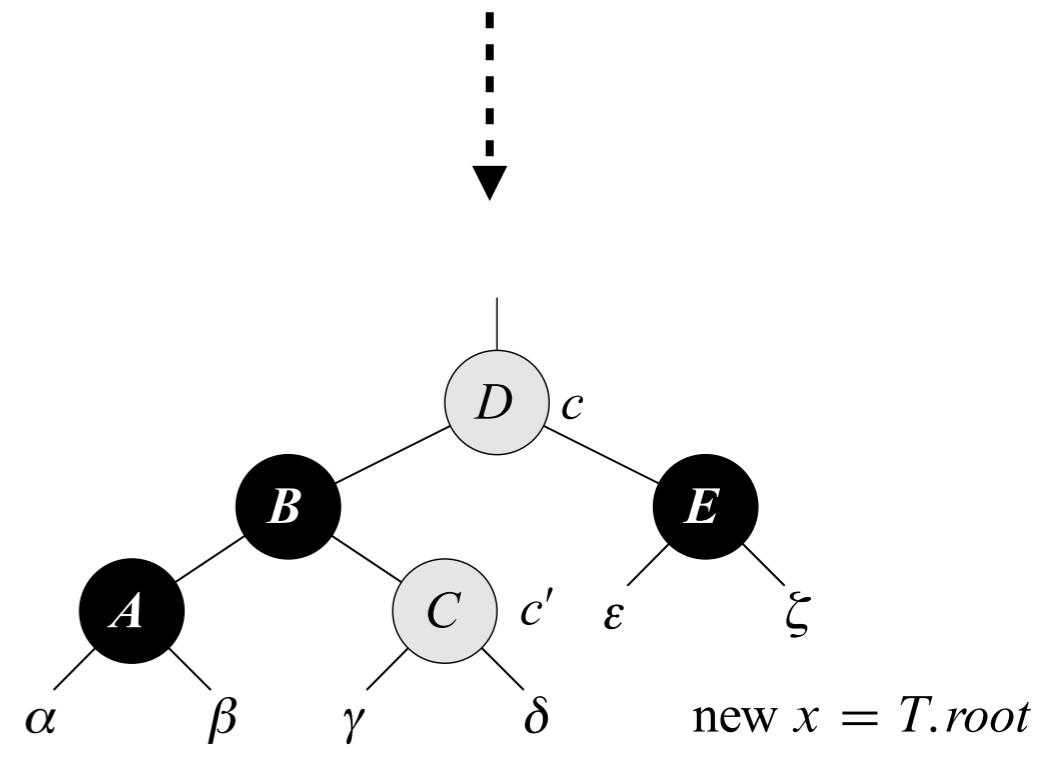
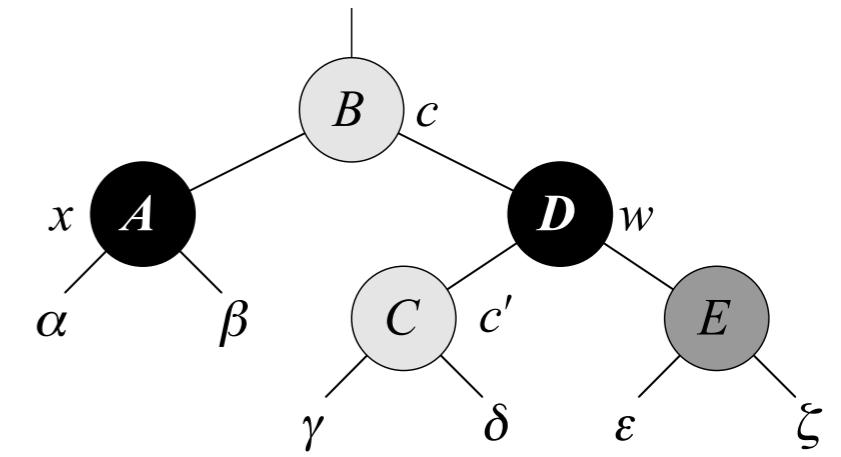
RB-Delete-Fixup: Case 4

RB-DELETE-FIXUP(T, x)

```

1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2    if  $x == x.p.left$ 
3       $w = x.p.right$ 
4      if  $w.color == \text{RED}$ 
5         $w.color = \text{BLACK}$ 
6         $x.p.color = \text{RED}$ 
7        LEFT-ROTATE( $T, x.p$ )
8         $w = x.p.right$ 
9      if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10         $w.color = \text{RED}$ 
11         $x = x.p$ 
12      else if  $w.right.color == \text{BLACK}$ 
13         $w.left.color = \text{BLACK}$ 
14         $w.color = \text{RED}$ 
15        RIGHT-ROTATE( $T, w$ )
16         $w = x.p.right$ 
17         $w.color = x.p.color$ 
18         $x.p.color = \text{BLACK}$ 
19         $w.right.color = \text{BLACK}$ 
20        LEFT-ROTATE( $T, x.p$ )
21         $x = T.root$ 
22      else (same as then clause with “right” and “left” exchanged)
23       $x.color = \text{BLACK}$ 
```

x 's sibling w is black, and
 w 's right child is red



RB-Delete: run-time

RB-DELETE(T, z)

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4     $x = z.\text{right}$ 
5    RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7     $x = z.\text{left}$ 
8    RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10    $y\text{-original-color} = y.\text{color}$ 
11    $x = y.\text{right}$ 
12   if  $y.p == z$ 
13      $x.p = y$ 
14   else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15      $y.\text{right} = z.\text{right}$ 
16      $y.\text{right}.p = y$ 
17   RB-TRANSPLANT( $T, z, y$ )
18    $y.\text{left} = z.\text{left}$ 
19    $y.\text{left}.p = y$ 
20    $y.\text{color} = z.\text{color}$ 
21   if  $y\text{-original-color} == \text{BLACK}$ 
22     RB-DELETE-FIXUP( $T, x$ )
```

$\Theta(h)$ as deletion in BST
+
 $\Theta(1)$

$\Theta(h)$. The number of loops
is bounded by h

RB Tres: Summary & Reflections

- As BSTs, Red-black tree are an effective data structure for dynamic sets
- They have in particular $h = O(\lg n)$ where h is the height of the BST
- **All the operations** we have seen **run in at most $O(\lg n)$ time**, where n is the number of elements in the tree
- Insertion and Deletion require 2 traverse of the tree:
 - One as for BST insertion and Deletion
 - One for fixing the RB property